

DOTS

Internet-Draft

Intended status: Standards Track

Expires: ~~January 28~~, April 1, 2021

M. Boucadair, Ed.

Orange

T. Reddy, Ed.

E. Doron

Radware Ltd.

M. Chen

CMCC

J. Shallow

~~July 27~~

September 28, 2020

McAfee

Distributed Denial-of-Service Open Threat Signaling (DOTS) Telemetry
draft-ietf-dots-telemetry-11

Abstract

This document aims to enrich DOTS signal channel protocol with various telemetry attributes allowing optimal Distributed Denial-of-Service attack mitigation. It specifies the normal traffic baseline and attack traffic telemetry attributes a DOTS client can convey to its DOTS server in the mitigation request, the mitigation status telemetry attributes a DOTS server can communicate to a DOTS client, and the mitigation efficacy telemetry attributes a DOTS client can communicate to a DOTS server. The telemetry attributes can assist the mitigator to choose the DDoS mitigation techniques and perform optimal DDoS attack mitigation.

Status of This Memo

This Internet-Draft is submitted in full conformance with the provisions of BCP 78 and BCP 79.

Internet-Drafts are working documents of the Internet Engineering Task Force (IETF). Note that other groups may also distribute working documents as Internet-Drafts. The list of current Internet-Drafts is at <https://datatracker.ietf.org/drafts/current/>.

Internet-Drafts are draft documents valid for a maximum of six months and may be updated, replaced, or obsoleted by other documents at any time. It is inappropriate to use Internet-Drafts as reference material or to cite them other than as "work in progress."

This Internet-Draft will expire on ~~January 28~~, April 1, 2021.

Copyright Notice

Copyright (c) 2020 IETF Trust and the persons identified as the document authors. All rights reserved.

This document is subject to BCP 78 and the IETF Trust's Legal Provisions Relating to IETF Documents (<https://trustee.ietf.org/license-info>) in effect on the date of publication of this document. Please review these documents carefully, as they describe your rights and restrictions with respect to this document. Code Components extracted from this document must include Simplified BSD License text as described in Section 4.e of the Trust Legal Provisions and are provided without warranty as described in the Simplified BSD License.

Table of Contents

1. Introduction	3
2. Terminology	5
3. DOTS Telemetry: Overview and Purpose	6
3.1. Need More Visibility	6
3.2. Enhanced Detection	7
3.3. Efficient Mitigation	9
4. Design Overview	9
4.1. Overview of Telemetry Operations	9
4.2. Generic Considerations	10
4.2.1. DOTS Client Identification	10
4.2.2. DOTS Gateways	10
4.2.3. Empty URI Paths	10
4.2.4. Controlling Configuration Data	10
4.3. Block-wise Transfer	11
4.4. DOTS Multi-homing Considerations	11
4.5. YANG Considerations	11
4.6. A Note About Examples	12
5. Telemetry Operation Paths	12
6. DOTS Telemetry Setup Configuration	13
6.1. Telemetry Configuration	14
6.1.1. Retrieve Current DOTS Telemetry Configuration	14
6.1.2. Convey DOTS Telemetry Configuration	16
6.1.3. Retrieve Installed DOTS Telemetry Configuration	20 19
6.1.4. Delete DOTS Telemetry Configuration	20

6.2.	Total Pipe Capacity	21	20
6.2.1.	Convey DOTS Client Domain Pipe Capacity	22	21
6.2.2.	Retrieve Installed DOTS Client Domain Pipe Capacity		27
6.2.3.	Delete Installed DOTS Client Domain Pipe Capacity		27
6.3.	Telemetry Baseline	28	27
6.3.1.	Convey DOTS Client Domain Baseline Information	31	30
6.3.2.	Retrieve Installed Normal Traffic Baseline	34	33
6.3.3.	Delete Installed Normal Traffic Baseline	34	33
6.4.	Reset Installed Telemetry Setup	34	33
6.5.	Conflict with Other DOTS Clients of the Same Domain	34	33
7.	DOTS Pre-or-Ongoing Mitigation Telemetry	35	34
7.1.	Pre-or-Ongoing-Mitigation DOTS Telemetry Attributes	37	36
7.1.1.	Target	38	37
7.1.2.	Total Traffic	39	38
7.1.3.	Total Attack Traffic	40	39
7.1.4.	Total Attack Connections	42	41
7.1.5.	Attack Details	45	44
7.2.	From DOTS Clients to DOTS Servers	51	50
7.3.	From DOTS Servers to DOTS Clients	54	53
8.	DOTS Telemetry Mitigation Status Update	59	58
8.1.	DOTS Clients to Servers Mitigation Efficacy DOTS Telemetry Attributes	59	58
8.2.	DOTS Servers to Clients Mitigation Status DOTS Telemetry Attributes	61	60
9.	Error Handling	65	64
10.	YANG Modules	65	64
10.1.	DOTS Signal Channel Telemetry YANG Module	65	64
10.2.	Vendor Attack Mapping Details YANG Module	93	92
11.	YANG/JSON Mapping Parameters to CBOR	96	95
12.	IANA Considerations	98	97
12.1.	DOTS Signal Channel CBOR Key Values	98	97
12.2.	DOTS Signal Channel Conflict Cause Codes	101	100
12.3.	DOTS Signal Telemetry YANG Module	101	100
13.	Security Considerations	102	101
13.1.	DOTS Signal Channel Telemetry	102	101
13.2.	Vendor Attack Mapping	103	102
14.	Contributors	104	103
15.	Acknowledgements	104	103
16.	References	104	103
16.1.	Normative References	104	103
16.2.	Informative References	106	105
	Authors' Addresses	107	106

1. Introduction

Distributed Denial of Service (DDoS) attacks have become more sophisticated. IT organizations and service providers are facing DDoS attacks that fall into two broad categories:

1. Network/Transport layer attacks target the victim's infrastructure. These attacks are not necessarily aimed at taking down the actual delivered services, but rather to eliminate various network elements (routers, switches, firewalls, transit links, and so on) from serving legitimate users traffic.

The main method of such attacks is to send a large volume or high packet per second (pps) of traffic toward the victim's infrastructure. Typically, attack volumes may vary from a few 100 Mbps to 100s of Gbps or even Tbps. Attacks are commonly carried out leveraging botnets and attack reflectors for amplification attacks such as NTP (Network Time Protocol), DNS (Domain Name System), SNMP (Simple Network Management Protocol), or SSDP (Simple Service Discovery Protocol).

2. Application layer attacks target various applications. Typical examples include attacks against HTTP/HTTPS, DNS, SIP (Session Initiation Protocol), or SMTP (Simple Mail Transfer Protocol). However, all applications with their port numbers open at network edges can be attractive attack targets.

Application layer attacks are considered more complex and hard to categorize, therefore harder to detect and mitigate efficiently.

To compound the problem, attackers also leverage multi-vectorized attacks. These attacks are assembled from dynamic attack vectors (Network/Application) and tactics. As such, multiple attack vectors formed by multiple attack types and volumes are launched simultaneously towards a victim. Multi-vector attacks are harder to detect and defend. Multiple and simultaneous mitigation techniques are needed to defeat such attack campaigns. It is also common for attackers to change attack vectors right after a successful mitigation, burdening their opponents with changing their defense methods.

The conclusion derived from these real scenarios is that modern attacks detection and mitigation are most certainly complicated and highly convoluted tasks. They demand a comprehensive knowledge of the attack attributes, the targeted normal behavior (including,

normal traffic patterns), as well as the attacker's on-going and past actions. Even more challenging, retrieving all the analytics needed for detecting these attacks is not simple to obtain with the industry's current capabilities.

The DOTS signal channel protocol ~~[I-D.boucadiar-dots-rfc8782-bis]~~ [I-D.ietf-dots-rfc8782-bis] is used to carry information about a network resource or a network (or a part thereof) that is under a DDoS attack. Such information is sent by a DOTS client to one or multiple DOTS servers so that appropriate mitigation actions are undertaken on traffic deemed suspicious. Various use cases are discussed in [I-D.ietf-dots-use-cases].

Typically, DOTS clients can be integrated within a DDoS attack detector, or network and security elements that have been actively engaged with ongoing attacks. The DOTS client mitigation environment determines that it is no longer possible or practical for it to handle these attacks. This can be due to a lack of resources or security capabilities, as derived from the complexities and the intensity of these attacks. In this circumstance, the DOTS client has invaluable knowledge about the actual attacks that need to be handled by its DOTS server(s). By enabling the DOTS client to share this comprehensive knowledge of an ongoing attack under specific circumstances, the DOTS server can drastically increase its ability to accomplish successful mitigation. While the attack is being handled by the DOTS server associated mitigation resources, the DOTS server has the knowledge about the ongoing attack mitigation. The DOTS server can share this information with the DOTS client so that the client can better assess and evaluate the actual mitigation realized.

DOTS clients can send mitigation hints derived from attack details to DOTS servers, with the full understanding that the DOTS server may ignore mitigation hints, as described in [RFC8612] (Gen-004). Mitigation hints will be transmitted across the DOTS signal channel, as the data channel may not be functional during an attack. How a DOTS server is handling normal and attack traffic attributes, and mitigation hints is implementation specific.

Both DOTS clients and servers can benefit this information by presenting various information in relevant management, reporting, and portal systems.

This document defines DOTS telemetry attributes that can be conveyed by DOTS clients to DOTS servers, and vice versa. The DOTS telemetry attributes are not mandatory attributes of the DOTS signal channel protocol ~~[I-D.boucadiar-dots-rfc8782-bis]~~, [I-D.ietf-dots-rfc8782-bis]. Nevertheless, when DOTS telemetry attributes are available to a DOTS agent, and absent any policy, it can signal the attributes in order to optimize the overall mitigation service provisioned using DOTS.

2. Terminology

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "NOT RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in BCP 14 [RFC2119][RFC8174] when, and only when, they appear in all capitals, as shown here.

The reader should be familiar with the terms defined in [RFC8612].

"DOTS Telemetry" is defined as the collection of attributes that are used to characterize normal traffic baseline, attacks and their mitigation measures, and any related information that may help in enforcing countermeasures. The DOTS Telemetry is an optional set of attributes that can be signaled in the DOTS signal channel protocol.

Telemetry Setup Identifier (tsid) is an identifier that is generated by DOTS clients to uniquely identify DOTS telemetry setup configuration data.

Telemetry Identifier (tmid) is an identifier that is generated by DOTS clients to uniquely identify DOTS telemetry data that is communicated prior or during a mitigation.

The meaning of the symbols in YANG tree diagrams are defined in [RFC8340] and [RFC8791].

3. DOTS Telemetry: Overview and Purpose

Timely and effective signaling of up-to-date DDoS telemetry to all elements involved in the mitigation process is essential and improves the overall DDoS mitigation service effectiveness. Bi-directional feedback between DOTS agents is required for an increased awareness of each party, supporting superior and highly efficient attack mitigation service.

3.1. Need More Visibility

When signaling a mitigation request, it is most certainly beneficial for DOTS clients to signal to DOTS servers any knowledge regarding ongoing attacks. This can happen in cases where DOTS clients are asking DOTS servers for support in defending against attacks that they have already detected and/or mitigated.

If attacks are already detected and categorized within a DOTS client domain, the DOTS server, and its associated mitigation services, can proactively benefit this information and optimize the overall service delivery. It is important to note that DOTS client domains and DOTS server domains detection and mitigation approaches can be different, and can potentially outcome different results and attack classifications. The DDoS mitigation service treats the ongoing attack details received from DOTS clients as hints and cannot completely rely or trust the attack details conveyed by DOTS clients.

A basic requirement of security operation teams is to be aware and get visibility into the attacks they need to handle. The DOTS server security operation teams benefit from the DOTS telemetry, especially from the reports of ongoing attacks. Even if some mitigation can be automated, operational teams can use the DOTS telemetry to be prepared for attack mitigation and to assign the correct resources (operation staff, networking and mitigation) for the specific service. Similarly, security operation personnel at the DOTS client side ask for feedback about their requests for protection. Therefore, it is valuable for DOTS servers to share DOTS telemetry with DOTS clients.

Mutual sharing of information is thus crucial for "closing the mitigation loop" between DOTS clients and servers. For the server side team, it is important to realize that the same attacks that the DOTS server's mitigation resources are seeing are those that a DOTS client is asking to mitigate. For the DOTS client side team, it is important to realize that the DOTS clients receive the required service. For example, understanding that "I asked for mitigation of two attacks and my DOTS server detects and mitigates only one of them". Cases of inconsistency in attack classification between DOTS clients and servers can be highlighted, and maybe handled, using the DOTS telemetry attributes.

In addition, management and orchestration systems, at both DOTS client and server sides, can use DOTS telemetry as a feedback to automate various control and management activities derived from signaled telemetry information.

If the DOTS server's mitigation resources have the capabilities to facilitate the DOTS telemetry, the DOTS server adapts its protection strategy and activates the required countermeasures immediately (automation enabled) for the sake of optimized attack mitigation decisions and actions.

3.2. Enhanced Detection

DOTS telemetry can also be used to tune the DDoS mitigators with the correct state of an attack. During the last few years, DDoS attack detection technologies have evolved from threshold-based detection (that is, cases when all or specific parts of traffic cross a predefined threshold for a certain period of time is considered as an attack) to an "anomaly detection" approach. For the latter, it is required to maintain rigorous learning of "normal" behavior and where an "anomaly" (or an attack) is identified and categorized based on the knowledge about the normal behavior and a deviation from this normal behavior. Machine learning approaches are used such that the actual traffic thresholds are automatically calculated by learning the protected entity normal traffic behavior during idle time. The normal traffic characterization learned is referred to as the "normal traffic baseline". An attack is detected when the victim's actual traffic is deviating from this normal baseline.

In addition, subsequent activities toward mitigating an attack are much more challenging. The ability to distinguish legitimate traffic from attacker traffic on a per packet basis is complex. For example, a packet may look "legitimate" and no attack signature can be identified. The anomaly can be identified only after detailed statistical analysis. DDoS attack mitigators use the normal baseline during the mitigation of an attack to identify and categorize the expected appearance of a specific traffic pattern. Particularly, the mitigators use the normal baseline to recognize the "level of normality" needs to be achieved during the various mitigation process.

Normal baseline calculation is performed based on continuous learning of the normal behavior of the protected entities. The minimum learning period varies from hours to days and even weeks, depending on the protected application behavior. The baseline cannot be learned during active attacks because attack conditions do not characterize the protected entities' normal behavior.

If the DOTS client has calculated the normal baseline of its protected entities, signaling such information to the DOTS server along with the attack traffic levels is significantly valuable. The DOTS server benefits from this telemetry by tuning its mitigation resources with the DOTS client's normal baseline. The DOTS server mitigators use the baseline to familiarize themselves with the attack victim's normal behavior and target the baseline as the level of normality they need to achieve. Fed with this information, the overall mitigation performances is expected to be improved in terms of time to mitigate, accuracy, false-negative, and false-positive.

Mitigation of attacks without having certain knowledge of normal traffic can be inaccurate at best. This is especially true for recursive signaling (see Section 3.2.3 in [I-D.ietf-dots-use-cases]). In addition, the highly diverse types of use cases where DOTS clients are integrated also emphasize the need for knowledge of each DOTS client domain behavior. Consequently, common global thresholds for attack detection practically cannot be realized. Each DOTS client domain can have its own levels of traffic and normal behavior. Without facilitating normal baseline signaling, it may be very difficult for DOTS servers in some cases to detect and mitigate the attacks accurately:

It is important to emphasize that it is practically impossible for the DOTS server's mitigators to calculate the normal baseline in cases where they do not have any knowledge of the traffic beforehand.

In addition, baseline learning requires a period of time that cannot be afforded during active attack.

Of course, this information can be provided using out-of-band mechanisms or manual configuration at the risk to maintain inaccurate information as the network evolves and "normal" patterns change. The use of a dynamic and collaborative means between the DOTS client and server to identify and share key parameters for the sake of efficient DDoS protection is valuable.

3.3. Efficient Mitigation

During a high volume attack, DOTS client pipes can be totally saturated. DOTS clients ask their DOTS servers to handle the attack upstream so that DOTS client pipes return to a reasonable load level (normal pattern, ideally). At this point, it is essential to ensure that the mitigator does not overwhelm the DOTS client pipes by sending back "clean traffic", or what it believes is "clean". This can happen when the mitigator has not managed to detect and mitigate all the attacks launched towards the DOTS client domain.

In this case, it can be valuable to DOTS clients to signal to DOTS servers the "total pipe capacity", which is the level of traffic the DOTS client domain can absorb from its upstream network. Dynamic updates of the condition of pipes between DOTS agents while they are under a DDoS attack is essential (e.g., where multiple DOTS clients share the same physical connectivity pipes). It is important to note that the term "pipe" noted here does not necessarily represent physical pipe, but rather represents the maximum level of traffic that the DOTS client domain can receive. The DOTS server should activate other mechanisms to ensure it does not allow the DOTS client domain's pipes to be saturated unintentionally. The rate-limit action defined in [RFC8783] is a reasonable candidate to achieve this objective; the DOTS client can ask for the type(s) of traffic (such as ICMP, UDP, TCP port number 80) it prefers to limit. The rate-limit action can be controlled via the signal channel [I-D.ietf-dots-signal-filter-control] even when the pipe is overwhelmed.

4. Design Overview

4.1. Overview of Telemetry Operations

This document specifies an extension to the DOTS signal channel protocol. Considerations about how to establish, maintain, and make use of the DOTS signal channel are specified in ~~[I-D.boucadiar-dots-rfc8782-bis].~~
[I-D.ietf-dots-rfc8782-bis].

Once the DOTS signal channel is established, DOTS clients that support the DOTS telemetry extension proceed with the telemetry setup configuration (e.g., measurement interval, telemetry notification interface, pipe capacity, normal traffic baseline) as detailed in Section 6. DOTS agents can then include DOTS telemetry attributes using the DOTS signal channel (Section 7.1). Typically, a DOTS client can use separate messages to share with its DOTS server(s) a set of telemetry data bound to an ongoing mitigation (Section 7.2). Also, a DOTS client that is interested to receive telemetry notifications related to some of its resources follows the procedure defined in Section 7.3. The DOTS client can then decide to send a

mitigation request if the notified attack cannot be mitigated locally within the DOTS client domain.

Aggregate DOTS telemetry data can also be included in efficacy update (Section 8.1) or mitigation update (Section 8.2) messages.

4.2. Generic Considerations

4.2.1. DOTS Client Identification

Following the rules in Section ~~5.4.1~~ 4.4.1 of ~~[I-D.boucadair-dots-rfc8782-bis]~~, ~~[I-D.ietf-dots-rfc8782-bis]~~, a unique identifier is generated by a DOTS client to prevent request collisions ('cuid').

As a reminder, ~~[I-D.boucadair-dots-rfc8782-bis]~~ ~~[I-D.ietf-dots-rfc8782-bis]~~ forbids 'cuid' to be returned in a response message body.

4.2.2. DOTS Gateways

DOTS gateways may be located between DOTS clients and servers. The considerations elaborated in Section ~~5.4.1~~ 4.4.1 of ~~[I-D.boucadair-dots-rfc8782-bis]~~ ~~[I-D.ietf-dots-rfc8782-bis]~~ must be followed. In particular, 'cddid' attribute is used to unambiguously identify a DOTS client domain.

As a reminder, ~~[I-D.boucadair-dots-rfc8782-bis]~~ ~~[I-D.ietf-dots-rfc8782-bis]~~ forbids 'cddid' (if present) to be returned in a response message body.

4.2.3. Empty URI Paths

Uri-Path parameters and attributes with empty values MUST NOT be present in a request and render an entire message invalid.

4.2.4. Controlling Configuration Data

The DOTS server follows the same considerations discussed in Section of ~~5.5.3~~ 4.5.3 of ~~[I-D.boucadair-dots-rfc8782-bis]~~ ~~[I-D.ietf-dots-rfc8782-bis]~~ for managing DOTS telemetry configuration freshness and notification.

Likewise, a DOTS client may control the selection of configuration and non-configuration data nodes when sending a GET request by means of the 'c' Uri-Query option and following the procedure specified in Section of ~~5.4.2~~ 4.4.2 of ~~[I-D.boucadair-dots-rfc8782-bis]~~, ~~[I-D.ietf-dots-rfc8782-bis]~~. These considerations are not reiterated in the following sections.

4.3. Block-wise Transfer

DOTS clients can use block wise transfer [RFC7959] with the recommendation detailed in Section ~~5.4.2~~ 4.4.2 of ~~[I-D.boucadair-dots-rfc8782-bis]~~ ~~[I-D.ietf-dots-rfc8782-bis]~~ to control the size of a response when the data to be returned does not fit within a single datagram.

DOTS clients can also use CoAP Block1 Option in a PUT request (see Section 2.5 of [RFC7959]) to initiate large transfers, but these Block1 transfers will fail if the inbound "pipe" is running full, so consideration needs to be made to try to fit this PUT into a single transfer, or to separate out the PUT into several discrete PUTs where each of them fits into a single packet.

Block3 and Block 4 Options that are similar to the CoAP Block1 and Block2 Options, but enable faster transmissions of big blocks of data with less packet interchanges, are defined in ~~[I-D.bosh-core-new-block]~~, ~~[I-D.ietf-core-new-block]~~. DOTS implementations can consider the use of Block3 and Block 4 Options.

4.4. DOTS Multi-homing Considerations

Multi-homed DOTS clients are assumed to follow the recommendations in [I-D.ietf-dots-multihoming] to select which DOTS server to contact and which IP prefixes to include in a telemetry message to a given peer DOTS server. For example, if each upstream network exposes a DOTS server and the DOTS client maintains DOTS channels with all of them, only the information related to prefixes assigned by an upstream network to the DOTS client domain will be signaled via the DOTS channel established with the DOTS server of that upstream network.

Considerations related to whether (and how) a DOTS client gleans some telemetry information (e.g., attack details) it receives from a first DOTS server and share it with a second DOTS server are implementation and deployment specific.

4.5. YANG Considerations

Telemetry messages exchanged between DOTS agents are serialized using

Concise Binary Object Representation (CBOR) [RFC7049]. CBOR-encoded payloads are used to carry signal channel specific payload messages which convey request parameters and response information such as errors.

This document specifies a YANG module [RFC7950] for representing DOTS telemetry message types (Section 10.1). All parameters in the payload of the DOTS signal channel are mapped to CBOR types as specified in Section 11.

The DOTS telemetry module (Section 10.1) is not intended to be used via NETCONF/RESTCONF for DOTS server management purposes. It serves only to provide a data model and encoding following [RFC8791].

The DOTS telemetry module (Section 10.1) uses "enumerations" rather than "identities" to define units, samples, and intervals because otherwise the namespace identifier "ietf-dots-telemetry" must be included when a telemetry attribute is included (e.g., in a mitigation efficacy update). The use of "identities" is thus suboptimal from a message compactness standpoint.

4.6. A Note About Examples

Examples are provided for illustration purposes. The document does not aim to provide a comprehensive list of message examples.

The authoritative reference for validating telemetry messages is the YANG module (Section 10.1) and the mapping table established in Section 11.

5. Telemetry Operation Paths

As discussed in Section 5.2 4.2 of ~~[I-D.boucadair-dots-rfc8782-bis]~~, [I-D.ietf-dots-rfc8782-bis], each DOTS operation is indicated by a path suffix that indicates the intended operation. The operation path is appended to the path prefix to form the URI used with a CoAP request to perform the desired DOTS operation. The following telemetry path suffixes are defined (Table 1):

Operation	Operation Path	Details
Telemetry Setup	/tm-setup	Section 6
Telemetry	/tm	Section 7

Table 1: DOTS Telemetry Operations

Consequently, the "ietf-dots-telemetry" YANG module defined in Section 10.1 defines data structure to represent new DOTS message types called 'telemetry-setup' and 'telemetry'. The tree structure is shown in Figure 1. More details are provided in Sections 6 and 7 about the exact structure of 'telemetry-setup' and 'telemetry' message types.

```
structure dots-telemetry:
  +-- (telemetry-message-type)?
    +--:(telemetry-setup)
      |
      | ..
      | +-- telemetry* []
      |   ..
      |   +-- (setup-type)?
      |     +--:(telemetry-config)
      |       | ..
      |       +--:(pipe)
      |       | ..
      |       +--:(baseline)
      |       | ..
      +--:(telemetry)
      | ..
      | ..
```

Figure 1: New DOTS Message Types (YANG Tree Structure)

6. DOTS Telemetry Setup Configuration

In reference to Figure 1, a DOTS telemetry setup message MUST include only telemetry-related configuration parameters (Section 6.1) or information about DOTS client domain pipe capacity (Section 6.2) or telemetry traffic baseline (Section 6.3). As such, requests that include a mix of telemetry configuration, pipe capacity, or traffic baseline MUST be rejected by DOTS servers with a 4.00 (Bad Request).

A DOTS client can reset all installed DOTS telemetry setup configuration data following the considerations detailed in Section 6.4.

A DOTS server may detect conflicts when processing requests related to DOTS client domain pipe capacity or telemetry traffic baseline

with requests from other DOTS clients of the same DOTS client domain. More details are included in Section 6.5.

Telemetry setup configuration is bound to a DOTS client domain. DOTS servers MUST NOT expect DOTS clients to send regular requests to refresh the telemetry setup configuration. Any available telemetry setup configuration has a validity timeout of the DOTS association with a DOTS client domain. DOTS servers MUST NOT reset 'tsid' because a session failed with a DOTS client. DOTS clients update their telemetry setup configuration upon change of a parameter that may impact attack mitigation.

DOTS telemetry setup configuration request and response messages are marked as Confirmable messages (Section 2.1 of [RFC7252]).

6.1. Telemetry Configuration

A DOTS client can negotiate with its server(s) a set of telemetry configuration parameters to be used for telemetry. Such parameters include:

- o Percentile-related measurement parameters
- o Measurement units
- o Acceptable percentile values
- o Telemetry notification interval
- o Acceptable Server-originated telemetry

Section 11.3 of [RFC2330] includes more details about computing percentiles.

6.1.1. Retrieve Current DOTS Telemetry Configuration

A GET request is used to obtain acceptable and current telemetry configuration parameters on the DOTS server. This request may include a 'cuid' Uri-Path when the request is relayed by a DOTS gateway. An example of such request is depicted in Figure 2.

```
Header: GET (Code=0.01)
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "tm-setup"
Uri-Path: "cuid=dz6pHjaADkaFTbjr0JGBpw"
```

Figure 2: GET to Retrieve Current and Acceptable DOTS Telemetry Configuration

Upon receipt of such request, and assuming no error is encountered by processing the request, the DOTS server replies with a 2.05 (Content) response that conveys the current and telemetry parameters acceptable by the DOTS server. The tree structure of the response message body is provided in Figure 3. Note that the response also includes any pipe (Section 6.2) and baseline information (Section 6.3) maintained by the DOTS server for this DOTS client.

DOTS servers that support the capability of sending telemetry information to DOTS clients prior or during a mitigation (Section 8.2) sets 'server-originated-telemetry' under 'max-config-values' to 'true' ('false' is used otherwise). If 'server-originated-telemetry' is not present in a response, this is equivalent to receiving a request with 'server-originated-telemetry' set to 'false'.

```
structure dots-telemetry:
  +-- (telemetry-message-type)?
  +--:(telemetry-setup)
  | +-- (direction)?
  | | +--:(server-to-client-only)
  | | | +-- max-config-values
  | | | | +-- measurement-interval?          interval
  | | | | +-- measurement-sample?            sample
  | | | | +-- low-percentile?                 percentile
  | | | | +-- mid-percentile?                 percentile
  | | | | +-- high-percentile?                percentile
  | | | | +-- server-originated-telemetry?    boolean
  | | | | +-- telemetry-notify-interval?      uint32
  | | | +-- min-config-values
  | | | | +-- measurement-interval?          interval
  | | | | +-- measurement-sample?            sample
  | | | | +-- low-percentile?                 percentile
  | | | | +-- mid-percentile?                 percentile
  | | | | +-- high-percentile?                percentile
  | | | | +-- telemetry-notify-interval?      uint32
  | | | +-- supported-units
  | | | | +-- unit-config* [unit]
```



```

| | | +-- unit          unit-type
| | | +-- unit-status   boolean
| | | +-- query-type*   query-type
| +-- telemetry* []
| | +-- (direction)?
| | | +--:(server-to-client-only)
| | | +-- tsid?          uint32
| | +-- (setup-type)?
| | | +--:(telemetry-config)
| | | | +-- current-config
| | | | | +-- measurement-interval?   interval
| | | | | +-- measurement-sample?     sample
| | | | | +-- low-percentile?          percentile
| | | | | +-- mid-percentile?          percentile
| | | | | +-- high-percentile?         percentile
| | | | +-- unit-config* [unit]
| | | | | +-- unit          unit-type
| | | | | +-- unit-status   boolean
| | | | +-- server-originated-telemetry? boolean
| | | | +-- telemetry-notify-interval? uint32
| | +--:(pipe)
| | | ...
| | +--:(baseline)
| | | ...
+--:(telemetry)
...

```

Figure 3: Telemetry Configuration Tree Structure

When both 'min-config-values' and 'max-config-values' attributes are present, the values carried in 'max-config-values' attributes MUST be greater or equal to their counterpart in 'min-config-values' attributes.

6.1.2. Convey DOTS Telemetry Configuration

PUT request is used to convey the configuration parameters for the telemetry data (e.g., low, mid, or high percentile values). For example, a DOTS client may contact its DOTS server to change the default percentile values used as baseline for telemetry data. Figure 3 lists the attributes that can be set by a DOTS client in such PUT request. An example of a DOTS client that modifies all percentile reference values is shown in Figure 4.

```

Header: PUT (Code=0.03)
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "tm-setup"
Uri-Path: "cuid=dz6pHjaADkaFTbjr0JGBpw"
Uri-Path: "tsid=123"
Content-Format: "application/dots+cbor"

```

```

{
  "ietf-dots-telemetry:telemetry-setup": {
    "telemetry": [
      {
        "current-config": {
          "low-percentile": "5.00",
          "mid-percentile": "65.00",
          "high-percentile": "95.00"
        }
      }
    ]
  }
}

```

Figure 4: PUT to Convey the DOTS Telemetry Configuration

'cuid' is a mandatory Uri-Path parameter for PUT requests.

The following additional Uri-Path parameter is defined:

tsid: Telemetry Setup Identifier is an identifier for the DOTS telemetry setup configuration data represented as an integer. This identifier MUST be generated by DOTS clients. 'tsid' values MUST increase monotonically (when a new PUT is generated by a DOTS client to convey new configuration parameters for the telemetry).

The procedure specified in Section 5.4.1 of ~~[I-D.boucadair-dots-rfc8782-bis]~~ **[I-D.ietf-dots-rfc8782-bis]** MUST be followed for 'tsid' rollover.

This is a mandatory attribute. 'tsid' MUST follow 'cuid'.

'cuid' and 'tsid' MUST NOT appear in the PUT request message body.

At least one configurable attribute MUST be present in the PUT request.

The PUT request with a higher numeric 'tsid' value overrides the DOTS telemetry configuration data installed by a PUT request with a lower numeric 'tsid' value. To avoid maintaining a long list of 'tsid' requests for requests carrying telemetry configuration data from a DOTS client, the lower numeric 'tsid' MUST be automatically deleted and no longer be available at the DOTS server.

The DOTS server indicates the result of processing the PUT request using the following Response Codes:

- o If the request is missing a mandatory attribute, does not include 'cuid' or 'tsid' Uri-Path parameters, or contains one or more invalid or unknown parameters, 4.00 (Bad Request) MUST be returned in the response.
- o If the DOTS server does not find the 'tsid' parameter value conveyed in the PUT request in its configuration data and if the DOTS server has accepted the configuration parameters, then a 2.01 (Created) Response Code MUST be returned in the response.
- o If the DOTS server finds the 'tsid' parameter value conveyed in the PUT request in its configuration data and if the DOTS server has accepted the updated configuration parameters, 2.04 (Changed) MUST be returned in the response.
- o If any of the enclosed configurable attribute values are not acceptable to the DOTS server (Section 6.1.1), 4.22 (Unprocessable Entity) MUST be returned in the response.

The DOTS client may retry and send the PUT request with updated attribute values acceptable to the DOTS server.

By default, low percentile (10th percentile), mid percentile (50th percentile), high percentile (90th percentile), and peak (100th percentile) values are used to represent telemetry data. Nevertheless, a DOTS client can disable some percentile types (low, mid, high). In particular, setting 'low-percentile' to '0.00' indicates that the DOTS client is not interested in receiving low-percentiles. Likewise, setting 'mid-percentile' (or 'high-percentile') to the same value as 'low-percentile' (or 'mid-percentile') indicates that the DOTS client is not interested in receiving mid-percentiles (or high-percentiles). For example, a DOTS client can send the request depicted in Figure 5 to inform the server that it is interested in receiving only high-percentiles. This assumes that the client will only use that percentile type when sharing telemetry data with the server.

```
Header: PUT (Code=0.03)
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "tm-setup"
Uri-Path: "cuid=dz6pHjaADkaFTbjr0JGBpw"
Uri-Path: "tsid=569"
Content-Format: "application/dots+cbor"
```

```
{
  "ietf-dots-telemetry:telemetry-setup": {
    "telemetry": [
      {
        "current-config": {
          "low-percentile": "0.00",
          "mid-percentile": "0.00",
          "high-percentile": "95.00"
        }
      }
    ]
  }
}
```

Figure 5: PUT to Disable Low- and Mid-Percentiles

DOTS clients can also configure the unit type(s) to be used for traffic-related telemetry data. Typically, the supported unit types are: packets per second, bits per second, and bytes per second.

DOTS clients that are interested to receive pre or ongoing mitigation telemetry (pre-or-ongoing-mitigation) information from a DOTS server (Section 8.2) MUST set 'server-originated-telemetry' to 'true'. If 'server-originated-telemetry' is not present in a PUT request, this is equivalent to receiving a request with 'server-originated-telemetry' set to 'false'. An example of a request to enable pre-or-ongoing-mitigation telemetry from DOTS servers is shown in Figure 6.

```
Header: PUT (Code=0.03)
Uri-Path: ".well-known"
```

```

Uri-Path: "dots"
Uri-Path: "tm-setup"
Uri-Path: "cuid=dz6pHjaADkaFTbjr0JGBpw"
Uri-Path: "tsid=569"
Content-Format: "application/dots+cbor"

{
  "ietf-dots-telemetry:telemetry-setup": {
    "telemetry": [
      {
        "current-config": {
          "server-originated-telemetry": true
        }
      }
    ]
  }
}

```

Figure 6: PUT to Enable Pre-or-ongoing-mitigation Telemetry from the DOTS server

6.1.3. Retrieve Installed DOTS Telemetry Configuration

A DOTS client may issue a GET message with 'tsid' Uri-Path parameter to retrieve the current DOTS telemetry configuration. An example of such request is depicted in Figure 7.

```

Header: GET (Code=0.01)
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "tm-setup"
Uri-Path: "cuid=dz6pHjaADkaFTbjr0JGBpw"
Uri-Path: "tsid=123"

```

Figure 7: GET to Retrieve Current DOTS Telemetry Configuration

If the DOTS server does not find the 'tsid' Uri-Path value conveyed in the GET request in its configuration data for the requesting DOTS client, it MUST respond with a 4.04 (Not Found) error Response Code.

6.1.4. Delete DOTS Telemetry Configuration

A DELETE request is used to delete the installed DOTS telemetry configuration data (Figure 8). 'cuid' and 'tsid' are mandatory Uri-Path parameters for such DELETE requests.

```

Header: DELETE (Code=0.04)
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "tm-setup"
Uri-Path: "cuid=dz6pHjaADkaFTbjr0JGBpw"
Uri-Path: "tsid=123"

```

Figure 8: Delete Telemetry Configuration

The DOTS server resets the DOTS telemetry configuration back to the default values and acknowledges a DOTS client's request to remove the DOTS telemetry configuration using 2.02 (Deleted) Response Code. A 2.02 (Deleted) Response Code is returned even if the 'tsid' parameter value conveyed in the DELETE request does not exist in its configuration data before the request.

Section 6.4 discusses the procedure to reset all DOTS telemetry setup configuration.

6.2. Total Pipe Capacity

A DOTS client can communicate to the DOTS server(s) its DOTS client domain pipe information. The tree structure of the pipe information is shown in Figure 9.

```

structure dots-telemetry:
  +-- (telemetry-message-type)?
  +--:(telemetry-setup)
  |   ...
  |   +-- telemetry* []
  |   |   +-- (direction)?
  |   |   |   +--:(server-to-client-only)
  |   |   |   |   +-- tsid?                uint32
  |   |   +-- (setup-type)?
  |   |   +--:(telemetry-config)
  |   |   |   ...
  |   |   +--:(pipe)
  |   |   |   +-- total-pipe-capacity* [link-id unit]
  |   |   |   |   +-- link-id            nt:link-id
  |   |   |   |   +-- capacity          uint64
  |   |   |   |   +-- unit
  |   |   +--:(baseline)

```

```

|           ...
+--:(telemetry)
|           ...

```

Figure 9: Pipe Tree Structure

A DOTS client domain pipe is defined as a list of limits of (incoming) traffic volume ('total-pipe-capacity') that can be forwarded over ingress interconnection links of a DOTS client domain. Each of these links is identified with a 'link-id' [RFC8345].

The unit used by a DOTS client when conveying pipe information is captured in 'unit' attribute.

6.2.1. Convey DOTS Client Domain Pipe Capacity

Similar considerations to those specified in Section 6.1.2 are followed with one exception:

The relative order of two PUT requests carrying DOTS client domain pipe attributes from a DOTS client is determined by comparing their respective 'tsid' values. If such two requests have overlapping 'link-id' and 'unit', the PUT request with higher numeric 'tsid' value will override the request with a lower numeric 'tsid' value. The overlapped lower numeric 'tsid' MUST be automatically deleted and no longer be available.

DOTS clients SHOULD minimize the number of active 'tsids' used for pipe information. Typically, in order to avoid maintaining a long list of 'tsids' for pipe information, it is RECOMMENDED that DOTS clients include in any request to update information related to a given link the information of other links (already communicated using a lower 'tsid' value). Doing so, this update request will override these existing requests and hence optimize the number of 'tsid' request per DOTS client.

- o Note: This assumes that all link information can fit in one single message.

For example, a DOTS client managing a single homed domain (Figure 10) can send a PUT request (shown in Figure 11) to communicate the capacity of "link1" used to connect to its ISP.




Figure 10: Single Homed DOTS Client Domain

```

Header: PUT (Code=0.03)
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "tm-setup"
Uri-Path: "cuid=dz6pHjaADkaFTbjr0JGBpw"
Uri-Path: "tsid=457"
Content-Format: "application/dots+cbor"

```

```

{
  "ietf-dots-telemetry:telemetry-setup": {
    "telemetry": [
      {
        "total-pipe-capacity": [
          {
            "link-id": "link1",
            "capacity": "500",
            "unit": "megabit-ps"
          }
        ]
      }
    ]
  }
}

```

Figure 11: Example of a PUT Request to Convey Pipe Information (Single Homed)

DOTS clients may be instructed to signal a link aggregate instead of individual links. For example, a DOTS client that manages a DOTS client domain having two interconnection links with an upstream ISP (Figure 12) can send a PUT request (shown in Figure 13) to communicate the aggregate link capacity with its ISP. Signalling individual or aggregate link capacity is deployment specific.






Figure 12: DOTS Client Domain with Two Interconnection Links

```
Header: PUT (Code=0.03)
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "tm-setup"
Uri-Path: "cuid=dz6pHjaADkaFTbjr0JGBpw"
Uri-Path: "tsid=896"
Content-Format: "application/dots+cbor"
```

```
{
  "ietf-dots-telemetry:telemetry-setup": {
    "telemetry": [
      {
        "total-pipe-capacity": [
          {
            "link-id": "aggregate",
            "capacity": "700",
            "unit": "megabit-ps"
          }
        ]
      }
    ]
  }
}
```

Figure 13: Example of a PUT Request to Convey Pipe Information (Aggregated Link)

Now consider that the DOTS client domain was upgraded to connect to an additional ISP (e.g., ISP#B of Figure 14), the DOTS client can inform a third-party DOTS server (that is, not hosted with ISP#A and ISP#B domains) about this update by sending the PUT request depicted in Figure 15. This request also includes information related to "link1" even if that link is not upgraded. Upon receipt of this request, the DOTS server removes the request with 'tsid=457' and updates its configuration base to maintain two links (link#1 and link#2).




Figure 14: Multi-Homed DOTS Client Domain

```
Header: PUT (Code=0.03)
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "tm-setup"
Uri-Path: "cuid=dz6pHjaADkaFTbjr0JGBpw"
Uri-Path: "tsid=458"
Content-Format: "application/dots+cbor"
```

```
{
  "ietf-dots-telemetry:telemetry-setup": {
    "telemetry": [
      {
        "total-pipe-capacity": [
          {
            "link-id": "link1",
            "capacity": "500",
            "unit": "megabit-ps"
          },
          {
            "link-id": "link2",
            "capacity": "500",
            "unit": "megabit-ps"
          }
        ]
      }
    ]
  }
}
```

Figure 15: Example of a PUT Request to Convey Pipe Information (Multi-Homed)

A DOTS client can delete a link by sending a PUT request with the 'capacity' attribute set to "0" if other links are still active for the same DOTS client domain (see Section 6.2.3 for other delete cases). For example, if a DOTS client domain re-homes (that is, it changes its ISP), the DOTS client can inform its DOTS server about this update (e.g., from the network configuration in Figure 10 to the one shown in Figure 16) by sending the PUT request depicted in Figure 17. Upon receipt of this request, and assuming no error is encountered when processing the request, the DOTS server removes "link1" from its configuration bases for this DOTS client domain. Note that if the DOTS server receives a PUT request with a 'capacity' attribute set to "0" for all included links, it MUST reject the request with a 4.00 (Bad Request).




Figure 16: Multi-Homed DOTS Client Domain

```

Header: PUT (Code=0.03)
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "tm-setup"
Uri-Path: "cuid=dz6pHjaADkaFTbjr0JGBpw"
Uri-Path: "tsid=459"
Content-Format: "application/dots+cbor"

{
  "ietf-dots-telemetry:telemetry-setup": {
    "telemetry": [
      {
        "total-pipe-capacity": [
          {
            "link-id": "link1",
            "capacity": "0",
            "unit": "megabit-ps"
          },
          {
            "link-id": "link2",
            "capacity": "500",
            "unit": "megabit-ps"
          }
        ]
      }
    ]
  }
}
  
```

Figure 17: Example of a PUT Request to Convey Pipe Information (Multi-Homed)

6.2.2. Retrieve Installed DOTS Client Domain Pipe Capacity

A GET request with 'tsid' Uri-Path parameter is used to retrieve a specific installed DOTS client domain pipe related information. The same procedure as defined in Section 6.1.3 is followed.

To retrieve all pipe information bound to a DOTS client, the DOTS client proceeds as specified in Section 6.1.1.

6.2.3. Delete Installed DOTS Client Domain Pipe Capacity

A DELETE request is used to delete the installed DOTS client domain pipe related information. The same procedure as defined in Section 6.1.4 is followed.

6.3. Telemetry Baseline

A DOTS client can communicate to its DOTS server(s) its normal traffic baseline and connections capacity:

Total traffic normal baseline: The percentile values representing the total traffic normal baseline. It can be represented for a target using 'total-traffic-normal'.

The traffic normal per protocol ('total-traffic-normal-per-protocol') baseline is represented for a target and is transport-

protocol specific.

The traffic normal per port number ('total-traffic-normal-per-port') baseline is represented for each port number bound to a target.

If the DOTS client negotiated percentile values and units (Section 6.1), these negotiated values will be used instead of the default ones.

Total connections capacity: If the target is subjected to resource consuming DDoS attacks, the following optional attributes for the target per transport protocol are useful to detect resource consuming DDoS attacks:

- * The maximum number of simultaneous connections that are allowed to the target.
- * The maximum number of simultaneous connections that are allowed to the target per client.
- * The maximum number of simultaneous embryonic connections that are allowed to the target. The term "embryonic connection" refers to a connection whose connection handshake is not finished. Embryonic connection is only possible in connection-oriented transport protocols like TCP or SCTP.
- * The maximum number of simultaneous embryonic connections that are allowed to the target per client.
- * The maximum number of connections allowed per second to the target.
- * The maximum number of connections allowed per second to the target per client.
- * The maximum number of requests allowed per second to the target.
- * The maximum number of requests allowed per second to the target per client.
- * The maximum number of partial requests allowed per second to the target. Attacks relying upon partial requests create a connection with a target but do not send a complete request (e.g., HTTP request).
- * The maximum number of partial requests allowed per second to the target per client.

The aggregate per transport protocol is captured in 'total-connection-capacity', while port specific capabilities are represented using 'total-connection-capacity-per-port'.

The tree structure of the normal traffic baseline is shown in Figure 18.

```

structure dots-telemetry:
  +-- (telemetry-message-type)?
    +--:(telemetry-setup)
      | ...
      | +-- telemetry* []
      |   +-- (direction)?
      |   | +--:(server-to-client-only)
      |   | | +-- tsid?                               uint32
      |   | +-- (setup-type)?
      |   | +--:(telemetry-config)
      |   | | ...
      |   | +--:(pipe)
      |   | | ...
      |   | +--:(baseline)
      |   |   +-- baseline* [id]
      |   |   | +-- id
      |   |   | | uint32
      |   |   | +-- target-prefix*
      |   |   | | inet:ip-prefix
      |   |   | +-- target-port-range* [lower-port]
      |   |   | | +-- lower-port   inet:port-number
      |   |   | | +-- upper-port?  inet:port-number
      |   |   | +-- target-protocol*                               uint8
      |   |   | +-- target-fqdn*
      |   |   | | inet:domain-name
      |   |   | +-- target-uri*
      |   |   | | inet:uri
      |   |   | +-- alias-name*
      |   |   | | string
      |   |   | +-- total-traffic-normal* [unit]
      |   |   | | +-- unit                               unit

```

```

|         | +-- low-percentile-g?      yang:gauge64
|         | +-- mid-percentile-g?     yang:gauge64
|         | +-- high-percentile-g?    yang:gauge64
|         | +-- peak-g?               yang:gauge64
|         +-- total-traffic-normal-per-protocol*
|             [unit protocol]
|         | +-- protocol               uint8
|         | +-- unit                   unit
|         | +-- low-percentile-g?     yang:gauge64
|         | +-- mid-percentile-g?     yang:gauge64
|         | +-- high-percentile-g?    yang:gauge64
|         | +-- peak-g?               yang:gauge64
|         +-- total-traffic-normal-per-port* [unit port]
|             | +-- port               inet:port-number
|             | +-- unit               unit
|             | +-- low-percentile-g?  yang:gauge64
|             | +-- mid-percentile-g?  yang:gauge64
|             | +-- high-percentile-g? yang:gauge64
|             | +-- peak-g?            yang:gauge64
|         +-- total-connection-capacity* [protocol]
|             | +-- protocol           uint8
|             | +-- connection?        uint64
|             | +-- connection-client? uint64
|             | +-- embryonic?         uint64
|             | +-- embryonic-client?  uint64
|             | +-- connection-ps?     uint64
|             | +-- connection-client-ps? uint64
|             | +-- request-ps?        uint64
|             | +-- request-client-ps? uint64
|             | +-- partial-request-ps? uint64
|             | +-- partial-request-client-ps? uint64
|         +-- total-connection-capacity-per-port*
|             [protocol port]
|             | +-- port
|             | |   inet:port-number
|             | +-- protocol           uint8
|             | +-- connection?        uint64
|             | +-- connection-client? uint64
|             | +-- embryonic?         uint64
|             | +-- embryonic-client?  uint64
|             | +-- connection-ps?     uint64
|             | +-- connection-client-ps? uint64
|             | +-- request-ps?        uint64
|             | +-- request-client-ps? uint64
|             | +-- partial-request-ps? uint64
|             | +-- partial-request-client-ps? uint64
+--:(telemetry)
...

```

Figure 18: Telemetry Baseline Tree Structure

6.3.1. Convey DOTS Client Domain Baseline Information

Similar considerations to those specified in Section 6.1.2 are followed with one exception:

The relative order of two PUT requests carrying DOTS client domain baseline attributes from a DOTS client is determined by comparing their respective 'tsid' values. If such two requests have overlapping targets, the PUT request with higher numeric 'tsid' value will override the request with a lower numeric 'tsid' value. The overlapped lower numeric 'tsid' MUST be automatically deleted and no longer be available.

Two PUT requests from a DOTS client have overlapping targets if there is a common IP address, IP prefix, FQDN, URI, or alias-name. Also, two PUT requests from a DOTS client have overlapping targets if the addresses associated with the FQDN, URI, or alias are overlapping with each other or with 'target-prefix'.

DOTS clients SHOULD minimize the number of active 'tsids' used for baseline information. Typically, in order to avoid maintaining a long list of 'tsids' for baseline information, it is RECOMMENDED that DOTS clients include in a request to update information related to a given target, the information of other targets (already communicated using a lower 'tsid' value) (assuming this fits within one single datagram). This update request will override these existing requests and hence optimize the number of 'tsid' request per DOTS client.

If no target clause is included in the request, this is an indication that the baseline information applies for the DOTS client domain as a whole.

An example of a PUT request to convey the baseline information is shown in Figure 19.

```
Header: PUT (Code=0.03)
Uri-Path: ".well-known"
```



```

Uri-Path: "dots"
Uri-Path: "tm-setup"
Uri-Path: "cuid=dz6pHjaADkaFTbjr0JGBpw"
Uri-Path: "tsid=126"
Content-Format: "application/dots+cbor"

{
  "ietf-dots-telemetry:telemetry-setup": {
    "telemetry": [
      {
        "baseline": [
          {
            "id": 1,
            "target-prefix": [
              "2001:db8:6401::1/128",
              "2001:db8:6401::2/128"
            ],
            "total-traffic-normal": [
              {
                "unit": "megabit-ps",
                "peak-g": "60"
              }
            ]
          }
        ]
      }
    ]
  }
}

```

Figure 19: PUT to Convey the DOTS Traffic Baseline

The DOTS client may share protocol specific baseline information (e.g., TCP and UDP) as shown in Figure 19.

```

Header: PUT (Code=0.03)
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "tm-setup"
Uri-Path: "cuid=dz6pHjaADkaFTbjr0JGBpw"
Uri-Path: "tsid=128"
Content-Format: "application/dots+cbor"

{
  "ietf-dots-telemetry:telemetry-setup": {
    "telemetry": [
      {
        "baseline": [
          {
            "id": 1,
            "target-prefix": [
              "2001:db8:6401::1/128",
              "2001:db8:6401::2/128"
            ],
            "total-traffic-normal-per-protocol": [
              {
                "unit": "megabit-ps",
                "protocol": 6,
                "peak-g": "50"
              },
              {
                "unit": "megabit-ps",
                "protocol": 17,
                "peak-g": "10"
              }
            ]
          }
        ]
      }
    ]
  }
}

```

Figure 20: PUT to Convey the DOTS Traffic Baseline (2)

The normal traffic baseline information should be updated to reflect legitimate overloads (e.g., flash crowds) to prevent unnecessary mitigation.

6.3.2. Retrieve Installed Normal Traffic Baseline

A GET request with 'tsid' Uri-Path parameter is used to retrieve a specific installed DOTS client domain baseline traffic information. The same procedure as defined in Section 6.1.3 is followed.

To retrieve all baseline information bound to a DOTS client, the DOTS client proceeds as specified in Section 6.1.1.

6.3.3. Delete Installed Normal Traffic Baseline

A DELETE request is used to delete the installed DOTS client domain normal traffic baseline. The same procedure as defined in Section 6.1.4 is followed.

6.4. Reset Installed Telemetry Setup

Upon bootstrapping (or reboot or any other event that may alter the DOTS client setup), a DOTS client MAY send a DELETE request to set the telemetry parameters to default values. Such a request does not include any 'tsid'. An example of such request is depicted in Figure 21.

```
Header: DELETE (Code=0.04)
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "tm-setup"
Uri-Path: "cuid=dz6pHjaADkaFTbjr0JGBpw"
```

Figure 21: Delete Telemetry Configuration

6.5. Conflict with Other DOTS Clients of the Same Domain

A DOTS server may detect conflicts between requests to convey pipe and baseline information received from DOTS clients of the same DOTS client domain. 'conflict-information' is used to report the conflict to the DOTS client following similar conflict handling discussed in Section 5.4.1 4.4.1 of ~~{I-D.boucadiar-dots-rfc8782-bis}~~. [I-D.ietf-dots-rfc8782-bis]. The conflict cause can be set to one of these values:

```
1: Overlapping targets (Section 5.4.1 4.4.1 of
{I-D.boucadiar-dots-rfc8782-bis}.
[I-D.ietf-dots-rfc8782-bis]).
```

TBA: Overlapping pipe scope (see Section 12).

7. DOTS Pre-or-Ongoing Mitigation Telemetry

There are two broad types of DDoS attacks, one is bandwidth consuming attack, the other is target resource consuming attack. This section outlines the set of DOTS telemetry attributes (Section 7.1) that covers both the types of attacks. The objective of these attributes is to allow for the complete knowledge of attacks and the various particulars that can best characterize attacks.

The "ietf-dots-telemetry" YANG module (Section 10.1) defines the data structure of a new message type called 'telemetry'. The tree structure of the 'telemetry' message type is shown in Figure 24.

The pre-or-ongoing-mitigation telemetry attributes are indicated by the path suffix '/tm'. The '/tm' is appended to the path prefix to form the URI used with a CoAP request to signal the DOTS telemetry. Pre-or-ongoing-mitigation telemetry attributes specified in Section 7.1 can be signaled between DOTS agents.

Pre-or-ongoing-mitigation telemetry attributes may be sent by a DOTS client or a DOTS server.

DOTS agents SHOULD bind pre-or-ongoing-mitigation telemetry data with mitigation requests relying upon the target clause. In particular, a telemetry PUT request sent after a mitigation request may include a reference to that mitigation request ('mid-list') as shown in Figure 22. An example illustrating requests correlation by means of 'target-prefix' is shown in Figure 23.

When generating telemetry data to send to a peer, the DOTS agent MUST auto-scale so that appropriate unit(s) are used.

```
+-----+
|DOTS client|
+-----+
|
|=====Mitigation Request (mid)=====>|
|
|===== Telemetry (mid-list{mid})=====>|
|
|
```

Figure 22: Example of Request Correlation using 'mid'

```
+-----+
|DOTS client|
+-----+
|
|<===== Telemetry (target-prefix)=====|
|
|=====Mitigation Request (target-prefix)=====>|
|
|
```

Figure 23: Example of Request Correlation using Target Prefix

DOTS agents MUST NOT send pre-or-ongoing-mitigation telemetry notifications to the same peer more frequently than once every 'telemetry-notify-interval' (Section 6.1). If a telemetry notification is sent using a block-like transfer mechanism (e.g., ~~{I-D.bosh-core-new-block}}~~, ~~[I-D.ietf-core-new-block]~~), this rate limit policy MUST NOT consider these individual blocks as separate notifications, but as a single notification.

DOTS pre-or-ongoing-mitigation telemetry request and response messages MUST be marked as Non-Confirmable messages (Section 2.1 of [RFC7252]).

```

structure dots-telemetry:
  +-- (telemetry-message-type)?
  +--:(telemetry-setup)
  |   ...
  |   +-- telemetry* []
  |   |   +-- (direction)?
  |   |   |   +--:(server-to-client-only)
  |   |   |   |   +-- tsid?                               uint32
  |   |   +-- (setup-type)?
  |   |   |   +--:(telemetry-config)
  |   |   |   |   ...
  |   |   +--:(pipe)
  |   |   |   ...
  |   |   +--:(baseline)
  |   |   |   ...
  +--:(telemetry)
  |   +-- pre-or-ongoing-mitigation* []
  |   |   +-- (direction)?
  |   |   |   +--:(server-to-client-only)
  |   |   |   |   +-- tmid?                               uint32
  |   |   +-- target
  |   |   |   ...
  |   |   +-- total-traffic* [unit]
  |   |   |   ...
  |   |   +-- total-traffic-protocol* [unit protocol]
  |   |   |   ...
  |   |   +-- total-traffic-port* [unit port]
  |   |   |   ...
  |   |   +-- total-attack-traffic* [unit]
  |   |   |   ...
  |   |   +-- total-attack-traffic-protocol* [unit protocol]
  |   |   |   ...
  |   |   +-- total-attack-traffic-port* [unit port]
  |   |   |   ...
  |   |   +-- total-attack-connection
  |   |   |   ...
  |   |   +-- total-attack-connection-port
  |   |   |   ...
  |   |   +-- attack-detail* [vendor-id attack-id]
  |   |   |   ...

```

Figure 24: Telemetry Message Type Tree Structure

7.1. Pre-or-Ongoing-Mitigation DOTS Telemetry Attributes

The description and motivation behind each attribute are presented in Section 3. DOTS telemetry attributes are optionally signaled and therefore MUST NOT be treated as mandatory fields in the DOTS signal channel protocol.

7.1.1. Target

A target resource (Figure 25) is identified using the attributes 'target-prefix', 'target-port-range', 'target-protocol', 'target-fqdn', 'target-uri', 'alias-name', or a pointer to a mitigation request ('mid-list').

```

+--:(telemetry)
  +-- pre-or-ongoing-mitigation* []
  |   +-- (direction)?
  |   |   +--:(server-to-client-only)
  |   |   |   +-- tmid?                               uint32
  |   +-- target
  |   |   +-- target-prefix*          inet:ip-prefix
  |   |   +-- target-port-range* [lower-port]
  |   |   |   +-- lower-port          inet:port-number
  |   |   |   +-- upper-port?         inet:port-number
  |   |   +-- target-protocol*        uint8
  |   |   +-- target-fqdn*            inet:domain-name
  |   |   +-- target-uri*              inet:uri
  |   |   +-- alias-name*              string
  |   |   +-- mid-list*                uint32

```

```

+-- total-traffic* [unit]
| ...
+-- total-traffic-protocol* [unit protocol]
| ...
+-- total-traffic-port* [unit port]
| ...
+-- total-attack-traffic* [unit]
| ...
+-- total-attack-traffic-protocol* [unit protocol]
| ...
+-- total-attack-traffic-port* [unit port]
| ...
+-- total-attack-connection
| ...
+-- total-attack-connection-port
| ...
+-- attack-detail* [vendor-id attack-id]
...

```

Figure 25: Target Tree Structure

At least one of the attributes 'target-prefix', 'target-fqdn', 'target-uri', 'alias-name', or 'mid-list' MUST be present in the target definition.

If the target is subjected to bandwidth consuming attack, the attributes representing the percentile values of the 'attack-id' attack traffic are included.

If the target is subjected to resource consuming DDoS attacks, the same attributes defined for Section 7.1.4 are applicable for representing the attack.

This is an optional subattribute.

7.1.2. Total Traffic

The 'total-traffic' attribute (Figure 26) conveys the percentile values of total traffic observed during a DDoS attack. More granular total traffic can be conveyed in 'total-traffic-protocol' and 'total-traffic-port'.

The 'total-traffic-protocol' represents the total traffic for a target and is transport-protocol specific.

The 'total-traffic-port' represents the total traffic for a target per port number.

```

+--:(telemetry)
+-- pre-or-ongoing-mitigation* []
+-- (direction)?
| +--:(server-to-client-only)
|   +-- tmid?                               uint32
+-- target
| ...
+-- total-traffic* [unit]
| +-- unit                                   unit
| +-- low-percentile-g?                     yang:gauge64
| +-- mid-percentile-g?                     yang:gauge64
| +-- high-percentile-g?                    yang:gauge64
| +-- peak-g?                              yang:gauge64
+-- total-traffic-protocol* [unit protocol]
| +-- protocol                             uint8
| +-- unit                                   unit
| +-- low-percentile-g?                     yang:gauge64
| +-- mid-percentile-g?                     yang:gauge64
| +-- high-percentile-g?                    yang:gauge64
| +-- peak-g?                              yang:gauge64
+-- total-traffic-port* [unit port]
| +-- port                                 inet:port-number
| +-- unit                                   unit
| +-- low-percentile-g?                     yang:gauge64
| +-- mid-percentile-g?                     yang:gauge64
| +-- high-percentile-g?                    yang:gauge64
| +-- peak-g?                              yang:gauge64
+-- total-attack-traffic* [unit]
| ...
+-- total-attack-traffic-protocol* [unit protocol]
| ...
+-- total-attack-traffic-port* [unit port]
| ...
+-- total-attack-connection
| ...
+-- total-attack-connection-port
| ...
+-- attack-detail* [vendor-id attack-id]
...

```

Figure 26: Total Traffic Tree Structure

7.1.3. Total Attack Traffic

The 'total-attack-traffic' attribute (Figure 27) conveys the total attack traffic identified by the DOTS client domain's DDoS Mitigation System (or DDoS Detector). More granular total traffic can be conveyed in 'total-attack-traffic-protocol' and 'total-attack-traffic-port'.

The 'total-attack-traffic-protocol' represents the total attack traffic for a target and is transport-protocol specific.

The 'total-attack-traffic-port' represents the total attack traffic for a target per port number.

```
+--:(telemetry)
+-- pre-or-ongoing-mitigation* []
+-- (direction)?
| +--:(server-to-client-only)
|   +-- tmid?                               uint32
+-- target
| ...
+-- total-traffic* [unit]
| ...
+-- total-traffic-protocol* [unit protocol]
| ...
+-- total-traffic-port* [unit port]
| ...
+-- total-attack-traffic* [unit]
| +-- protocol?                             uint8
| +-- unit                                  unit
| +-- low-percentile-g?                     yang:gauge64
| +-- mid-percentile-g?                    yang:gauge64
| +-- high-percentile-g?                   yang:gauge64
| +-- peak-g?                              yang:gauge64
+-- total-attack-traffic-protocol* [unit protocol]
| +-- protocol                             uint8
| +-- unit                                  unit
| +-- low-percentile-g?                    yang:gauge64
| +-- mid-percentile-g?                    yang:gauge64
| +-- high-percentile-g?                   yang:gauge64
| +-- peak-g?                             yang:gauge64
+-- total-attack-traffic-port* [unit port]
| +-- port                                 inet:port-number
| +-- unit                                  unit
| +-- low-percentile-g?                    yang:gauge64
| +-- mid-percentile-g?                    yang:gauge64
| +-- high-percentile-g?                   yang:gauge64
| +-- peak-g?                             yang:gauge64
+-- total-attack-connection
| ...
+-- total-attack-connection-port
| ...
+-- attack-detail* [vendor-id attack-id]
...
```

Figure 27: Total Attack Traffic Tree Structure

7.1.4. Total Attack Connections

If the target is subjected to resource consuming DDoS attack, the 'total-attack-connection' attribute is used to convey the percentile values of total attack connections. The following optional subattributes for the target per transport protocol are included to represent the attack characteristics:

- o The number of simultaneous attack connections to the target.
- o The number of simultaneous embryonic connections to the target.
- o The number of attack connections per second to the target.
- o The number of attack requests to the target.

The total attack connections per port number is represented using 'total-attack-connection-port' attribute.

```
+--:(telemetry)
+-- pre-or-ongoing-mitigation* []
+-- (direction)?
| +--:(server-to-client-only)
|   +-- tmid?                               uint32
+-- target
| ...
+-- total-traffic* [unit]
| ...
+-- total-traffic-protocol* [unit protocol]
| ...
+-- total-traffic-port* [unit port]
| ...
```

```

+-- total-attack-traffic* [unit]
| ...
+-- total-attack-traffic-protocol* [unit protocol]
| ...
+-- total-attack-traffic-port* [unit port]
| ...
+-- total-attack-connection
| +-- low-percentile-1* [protocol]
| | +-- protocol                uint8
| | +-- connection?            yang:gauge64
| | +-- embryonic?             yang:gauge64
| | +-- connection-ps?        yang:gauge64
| | +-- request-ps?           yang:gauge64
| | +-- partial-request-ps?    yang:gauge64
| | +-- mid-percentile-1* [protocol]
| | | +-- protocol                uint8
| | | +-- connection?            yang:gauge64
| | | +-- embryonic?             yang:gauge64
| | | +-- connection-ps?        yang:gauge64
| | | +-- request-ps?           yang:gauge64
| | | +-- partial-request-ps?    yang:gauge64
| | +-- high-percentile-1* [protocol]
| | | +-- protocol                uint8
| | | +-- connection?            yang:gauge64
| | | +-- embryonic?             yang:gauge64
| | | +-- connection-ps?        yang:gauge64
| | | +-- request-ps?           yang:gauge64
| | | +-- partial-request-ps?    yang:gauge64
| | +-- peak-1* [protocol]
| | | +-- protocol                uint8
| | | +-- connection?            yang:gauge64
| | | +-- embryonic?             yang:gauge64
| | | +-- connection-ps?        yang:gauge64
| | | +-- request-ps?           yang:gauge64
| | | +-- partial-request-ps?    yang:gauge64
| +-- total-attack-connection-port
| | +-- low-percentile-1* [protocol port]
| | | +-- port                    inet:port-number
| | | +-- protocol                uint8
| | | +-- connection?            yang:gauge64
| | | +-- embryonic?             yang:gauge64
| | | +-- connection-ps?        yang:gauge64
| | | +-- request-ps?           yang:gauge64
| | | +-- partial-request-ps?    yang:gauge64
| | | +-- mid-percentile-1* [protocol port]
| | | | +-- port                    inet:port-number
| | | | +-- protocol                uint8
| | | | +-- connection?            yang:gauge64
| | | | +-- embryonic?             yang:gauge64
| | | | +-- connection-ps?        yang:gauge64
| | | | +-- request-ps?           yang:gauge64
| | | | +-- partial-request-ps?    yang:gauge64
| | | +-- high-percentile-1* [protocol port]
| | | | +-- port                    inet:port-number
| | | | +-- protocol                uint8
| | | | +-- connection?            yang:gauge64
| | | | +-- embryonic?             yang:gauge64
| | | | +-- connection-ps?        yang:gauge64
| | | | +-- request-ps?           yang:gauge64
| | | | +-- partial-request-ps?    yang:gauge64
| | | +-- peak-1* [protocol port]
| | | | +-- port                    inet:port-number
| | | | +-- protocol                uint8
| | | | +-- connection?            yang:gauge64
| | | | +-- embryonic?             yang:gauge64
| | | | +-- connection-ps?        yang:gauge64
| | | | +-- request-ps?           yang:gauge64
| | | | +-- partial-request-ps?    yang:gauge64
| +-- attack-detail* [vendor-id attack-id]
| ...

```

Figure 28: Total Attack Connections Tree Structure

7.1.5. Attack Details

This attribute (Figure 29) is used to signal a set of details characterizing an attack. The following subattributes describing the ongoing attack can be signal as attack details.

vendor-id: Vendor ID is a security vendor's Enterprise Number as registered with IANA [Enterprise-Numbers]. It is a four-byte integer value.

attack-id: Unique identifier assigned for the attack.

attack-description: Textual representation of the attack description. Natural Language Processing techniques (e.g., word embedding) can possibly be used to map the attack description to

an attack type. Textual representation of attack solves two problems: (a) avoids the need to create mapping tables manually between vendors and (b) avoids the need to standardize attack types which keep evolving.

attack-severity: Attack severity level. This attribute takes one of the values defined in Section 3.12.2 of [RFC7970].

start-time: The time the attack started. The attack's start time is expressed in seconds relative to 1970-01-01T00:00Z in UTC time (Section 2.4.1 of [RFC7049]). The CBOR encoding is modified so that the leading tag 1 (epoch-based date/time) MUST be omitted.

end-time: The time the attack ended. The attack end time is expressed in seconds relative to 1970-01-01T00:00Z in UTC time (Section 2.4.1 of [RFC7049]). The CBOR encoding is modified so that the leading tag 1 (epoch-based date/time) MUST be omitted.

source-count: A count of sources involved in the attack targeting the victim.

top-talkers: A list of top talkers among attack sources. The top talkers are represented using the 'source-prefix'.

'spoofed-status' indicates whether a top talker is a spoofed IP address (e.g., reflection attacks) or not.

If the target is subjected to a bandwidth consuming attack, the attack traffic from each of the top talkers is included ('total-attack-traffic', Section 7.1.3).

If the target is subjected to a resource consuming DDoS attack, the same attributes defined in Section 7.1.4 are applicable for representing the attack per talker.

```
+--:(telemetry)
  +-- pre-or-ongoing-mitigation* []
    +-- (direction)?
      | +--:(server-to-client-only)
      |   +-- tmid?                               uint32
    +-- target
      | ...
    +-- total-traffic* [unit]
      | ...
    +-- total-traffic-protocol* [unit protocol]
      | ...
    +-- total-traffic-port* [unit port]
      | ...
    +-- total-attack-traffic* [unit]
      | ...
    +-- total-attack-traffic-protocol* [unit protocol]
      | ...
    +-- total-attack-traffic-port* [unit port]
      | ...
    +-- total-attack-connection
      | ...
    +-- total-attack-connection-port
      | ...
    +-- attack-detail* [vendor-id attack-id]
      +-- vendor-id                               uint32
      +-- attack-id                               uint32
      +-- attack-description?                      string
      +-- attack-severity?                        attack-severity
      +-- start-time?                             uint64
      +-- end-time?                               uint64
      +-- source-count
        | +-- low-percentile-g?                    yang:gauge64
        | +-- mid-percentile-g?                   yang:gauge64
        | +-- high-percentile-g?                  yang:gauge64
        | +-- peak-g?                             yang:gauge64
      +-- top-talker
        +-- talker* [source-prefix]
          +-- spoofed-status?                      boolean
          +-- source-prefix                        inet:ip-prefix
          +-- source-port-range* [lower-port]
            | +-- lower-port                      inet:port-number
            | +-- upper-port?                     inet:port-number
          +-- source-icmp-type-range* [lower-type]
            | +-- lower-type                      uint8
            | +-- upper-type?                     uint8
          +-- total-attack-traffic* [unit]
            | +-- unit                            unit
            | +-- low-percentile-g?                yang:gauge64
            | +-- mid-percentile-g?                yang:gauge64
            | +-- high-percentile-g?               yang:gauge64
            | +-- peak-g?                         yang:gauge64
          +-- total-attack-connection
            +-- low-percentile-l* [protocol]
```

```

| +-- protocol                uint8
| +-- connection?            yang:gauge64
| +-- embryonic?             yang:gauge64
| +-- connection-ps?         yang:gauge64
| +-- request-ps?            yang:gauge64
| +-- partial-request-ps?    yang:gauge64
+-- mid-percentile-1* [protocol]
| +-- protocol                uint8
| +-- connection?            yang:gauge64
| +-- embryonic?             yang:gauge64
| +-- connection-ps?         yang:gauge64
| +-- request-ps?            yang:gauge64
| +-- partial-request-ps?    yang:gauge64
+-- high-percentile-1* [protocol]
| +-- protocol                uint8
| +-- connection?            yang:gauge64
| +-- embryonic?             yang:gauge64
| +-- connection-ps?         yang:gauge64
| +-- request-ps?            yang:gauge64
| +-- partial-request-ps?    yang:gauge64
+-- peak-1* [protocol]
  +-- protocol                uint8
  +-- connection?            yang:gauge64
  +-- embryonic?             yang:gauge64
  +-- connection-ps?         yang:gauge64
  +-- request-ps?            yang:gauge64
  +-- partial-request-ps?    yang:gauge64

```

Figure 29: Attack Detail Tree Structure

In order to optimize the size of telemetry data conveyed over the DOTS signal channel, DOTS agents MAY use the DOTS data channel [RFC8783] to exchange vendor specific attack mapping details (that is, {vendor identifier, attack identifier} ==> attack description). As such, DOTS agents do not have to convey systematically an attack description in their telemetry messages over the DOTS signal channel.

Multiple mappings for different vendor identifiers may be used; the DOTS agent transmitting telemetry information can elect to use one or more vendor mappings even in the same telemetry message.

Note: It is possible that a DOTS server is making use of multiple DOTS mitigators; each from a different vendor. How telemetry information and vendor mappings are exchanged between DOTS servers and DOTS mitigators is outside the scope of this document.

DOTS clients and servers may be provided with mappings from different vendors and so have their own different sets of vendor attack mappings. A DOTS agent MUST accept receipt of telemetry data with a vendor identifier that is different to the one it uses to transmit telemetry data. Furthermore, it is possible that the DOTS client and DOTS server are provided by the same vendor, but the vendor mapping tables are at different revisions. The DOTS client SHOULD transmit telemetry information using the vendor mapping(s) that it provided to the DOTS server and the DOTS server SHOULD use the vendor mappings(s) provided to the DOTS client when transmitting telemetry data to peer DOTS agent.

The "ietf-dots-mapping" YANG module defined in Section 10.2 augments the "ietf-dots-data-channel" [RFC8783]. The tree structure of this module is shown in Figure 30.

```

module: ietf-dots-mapping
  augment /ietf-data:dots-data/ietf-data:dots-client /data-channel:dots-data/data-channel:dots-client:
    +--rw vendor-mapping {dots-telemetry}?
      +--rw vendor* [vendor-id]
        +--rw vendor-id          uint32
        +--rw vendor-name?       string
        +--rw last-updated        uint64
        +--rw attack-mapping* [attack-id]
          +--rw attack-id        uint32
          +--rw attack-description string
  augment /ietf-data:dots-data/ietf-data:capabilities /data-channel:dots-data/data-channel:capabilities:
    +--ro vendor-mapping-enabled? boolean {dots-telemetry}?
  augment /ietf-data:dots-data /data-channel:dots-data:
    +--ro vendor-mapping {dots-telemetry}?
      +--ro vendor* [vendor-id]
        +--ro vendor-id          uint32
        +--ro vendor-name?       string
        +--ro last-updated        uint64
        +--ro attack-mapping* [attack-id]
          +--ro attack-id        uint32
          +--ro attack-description string

```

Figure 30: Vendor Attack Mapping Tree Structure

A DOTS client sends a GET request to retrieve the capabilities supported by a DOTS server as per Section 7.1 of [RFC8783]. This

request is meant to assess whether vendor attack mapping details feature is supported by the server (i.e., check the value of 'vendor-mapping-enabled').

If 'vendor-mapping-enabled' is set to 'true', A DOTS client MAY send a GET request to retrieve the DOTS server's vendor attack mapping details. An example of such GET request is shown in Figure 31.

```
GET /restconf/data/ietf-dots-data-channel:dots-data\
    /ietf-dots-mapping:vendor-mapping HTTP/1.1
Host: example.com
Accept: application/yang-data+json
```

Figure 31: GET to Retrieve the Vendor Attack Mappings of a DOTS Server

A DOTS client MAY retrieve only the list of vendors supported by the DOTS server. It does so by setting the "depth" parameter (Section 4.8.2 of [RFC8040]) to "3" in the GET request as shown in Figure 32. An example of a response body received from the DOTS server as a response to such request is illustrated in Figure 33.

```
GET /restconf/data/ietf-dots-data-channel:dots-data\
    /ietf-dots-mapping:vendor-mapping?depth=3 HTTP/1.1
Host: example.com
Accept: application/yang-data+json
```

Figure 32: GET to Retrieve the Vendors List used by a DOTS Server

```
{
  "ietf-dots-mapping:vendor-mapping": {
    "vendor": [
      {
        "vendor-id": 1234,
        "vendor-name": "mitigator-s",
        "last-updated": "1576856561",
        "attack-mapping": []
      }
    ]
  }
}
```

Figure 33: Response to a GET to Retrieve the Vendors List used by a DOTS Server

The DOTS client reiterates the above procedure regularly (e.g., once a week) to update the DOTS server's vendor attack mapping details.

If the DOTS client concludes that the DOTS server does not have any reference to the specific vendor attack mapping details, the DOTS client uses a POST request to install its vendor attack mapping details. An example of such POST request is depicted in Figure 34.

```
POST /restconf/data/ietf-dots-data-channel:dots-data\
    /dots-client=dz6pHjaADkaFTbjr0JGBpw HTTP/1.1
Host: example.com
Content-Type: application/yang-data+json
```

```
{
  "ietf-dots-mapping:vendor-mapping": {
    "vendor": [
      {
        "vendor-id": 345,
        "vendor-name": "mitigator-c",
        "last-updated": "1576812345",
        "attack-mapping": [
          {
            "attack-id": 1,
            "attack-description":
              "Include a description of this attack"
          },
          {
            "attack-id": 2,
            "attack-description":
              "Again, include a description of the attack"
          }
        ]
      }
    ]
  }
}
```

Figure 34: POST to Install Vendor Attack Mapping Details

The DOTS server indicates the result of processing the POST request using the status-line. Concretely, "201 Created" status-line MUST be returned in the response if the DOTS server has accepted the vendor attack mapping details. If the request is missing a mandatory

attribute or contains an invalid or unknown parameter, "400 Bad Request" status-line MUST be returned by the DOTS server in the response. The error-tag is set to "missing-attribute", "invalid-value", or "unknown-element" as a function of the encountered error.

If the request is received via a server-domain DOTS gateway, but the DOTS server does not maintain a 'cdid' for this 'cuid' while a 'cdid' is expected to be supplied, the DOTS server MUST reply with "403 Forbidden" status-line and the error-tag "access-denied". Upon receipt of this message, the DOTS client MUST register (Section 5.1 of [RFC8783]).

The DOTS client uses the PUT request to modify its vendor attack mapping details maintained by the DOTS server (e.g., add a new mapping).

A DOTS client uses a GET request to retrieve its vendor attack mapping details as maintained by the DOTS server (Figure 35).

```
GET /restconf/data/ietf-dots-data-channel:dots-data\
    /dots-client=dz6pHjaADkaFTbjr0JGBpw\
    /ietf-dots-mapping:vendor-mapping?\
    content=all HTTP/1.1
Host: example.com
Accept: application/yang-data+json
```

Figure 35: GET to Retrieve Installed Vendor Attack Mapping Details

When conveying attack details in DOTS telemetry messages (Sections 7.2, 7.3, and 8), DOTS agents MUST NOT include 'attack-description' attribute except if the corresponding attack mapping details were not shared with the peer DOTS agent.

7.2. From DOTS Clients to DOTS Servers

DOTS clients uses PUT request to signal pre-or-ongoing-mitigation telemetry to DOTS servers. An example of such request is shown in Figure 36.

```
Header: PUT (Code=0.03)
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "tm"
Uri-Path: "cuid=dz6pHjaADkaFTbjr0JGBpw"
Uri-Path: "tmid=123"
Content-Format: "application/dots+cbor"
```

```
{
  "ietf-dots-telemetry:telemetry": {
    "pre-or-ongoing-mitigation": [
      {
        "target": {
          "target-prefix": [
            "2001:db8::1/128"
          ]
        },
        "total-attack-traffic-protocol": [
          {
            "protocol": 17,
            "unit": "megabit-ps",
            "mid-percentile-g": "900"
          }
        ],
        "attack-detail": [
          {
            "vendor-id": 1234,
            "attack-id": 77,
            "start-time": "1957811234",
            "attack-severity": "high"
          }
        ]
      }
    ]
  }
}
```

Figure 36: PUT to Send Pre-or-Ongoing-Mitigation Telemetry

'cuid' is a mandatory Uri-Path parameter for PUT requests.

The following additional Uri-Path parameter is defined:

tmid: Telemetry Identifier is an identifier for the DOTS pre-or-ongoing-mitigation telemetry data represented as an integer. This identifier MUST be generated by DOTS clients. 'tmid' values MUST increase monotonically (when a new PUT is generated by a DOTS client to convey pre-or-ongoing-mitigation telemetry).

The procedure specified in Section ~~5.4.1~~ **4.4.1** of ~~[I-D.boucadiar-dots-rfc8782-bis]~~ **[I-D.ietf-dots-rfc8782-bis]** MUST be followed for 'tmid' rollover.

This is a mandatory attribute. 'tmid' MUST follow 'cuid'.

'cuid' and 'tmid' MUST NOT appear in the PUT request message body.

At least 'target' attribute and another pre-or-ongoing-mitigation attributes (Section 7.1) MUST be present in the PUT request. If only the 'target' attribute is present, this request is handled as per Section 7.3.

The relative order of two PUT requests carrying DOTS pre-or-ongoing-mitigation telemetry from a DOTS client is determined by comparing their respective 'tmid' values. If such two requests have overlapping 'target', the PUT request with higher numeric 'tmid' value will override the request with a lower numeric 'tmid' value. The overlapped lower numeric 'tmid' MUST be automatically deleted and no longer be available.

The DOTS server indicates the result of processing a PUT request using CoAP Response Codes. In particular, the 2.04 (Changed) Response Code is returned if the DOTS server has accepted the pre-or-ongoing-mitigation telemetry. The 5.03 (Service Unavailable) Response Code is returned if the DOTS server has erred. 5.03 uses Max-Age Option to indicate the number of seconds after which to retry.

How long a DOTS server maintains a 'tmid' as active or logs the enclosed telemetry information is implementation specific. Note that if a 'tmid' is still active, then logging details are updated by the DOTS server as a function of the updates received from the peer DOTS client.

A DOTS client that lost the state of its active 'tmids' or has to set 'tmid' back to zero (e.g., crash or restart) MUST send a GET request to the DOTS server to retrieve the list of active 'tmid'. The DOTS client may then delete 'tmids' that should not be active anymore (Figure 37). Sending a DELETE with no 'tmid' indicates that all 'tmids' must be deactivated (Figure 38).

```
Header: DELETE (Code=0.04)
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "tm"
Uri-Path: "cuid=dz6pHjaADkaFTbjr0JGBpw"
Uri-Path: "tmid=123"
```

Figure 37: Delete a Pre-or-Ongoing-Mitigation Telemetry

```
Header: DELETE (Code=0.04)
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "tm"
Uri-Path: "cuid=dz6pHjaADkaFTbjr0JGBpw"
```

Figure 38: Delete All Pre-or-Ongoing-Mitigation Telemetry

7.3. From DOTS Servers to DOTS Clients

The pre-or-ongoing-mitigation (attack details, in particular) can also be signaled from DOTS servers to DOTS clients. For example, the DOTS server co-located with a DDoS detector collects monitoring information from the target network, identifies DDoS attack using statistical analysis or deep learning techniques, and signals the attack details to the DOTS client.

The DOTS client can use the attack details to decide whether to trigger a DOTS mitigation request or not. Furthermore, the security operation personnel at the DOTS client domain can use the attack details to determine the protection strategy and select the appropriate DOTS server for mitigating the attack.

In order to receive pre-or-ongoing-mitigation telemetry notifications from a DOTS server, a DOTS client MUST send a PUT (followed by a GET) with the target filter. An example of such PUT request is shown in Figure 39. In order to avoid maintaining a long list of such requests, it is RECOMMENDED that DOTS clients include all targets in the same request. DOTS servers may be instructed to restrict the number of pre-or-ongoing-mitigation requests per DOTS client domain. This request MUST be maintained active by the DOTS server until a delete request is received from the same DOTS client to clear this pre-or-ongoing-mitigation telemetry.

The relative order of two PUT requests carrying DOTS pre-or-ongoing-mitigation telemetry from a DOTS client is determined by comparing

their respective 'tmid' values. If such two requests have overlapping 'target', the PUT request with higher numeric 'tmid' value will override the request with a lower numeric 'tmid' value.

The overlapped lower numeric 'tmid' MUST be automatically deleted and no longer be available.

```
Header: PUT (Code=0.03)
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "tm"
Uri-Path: "cuid=dz6pHjaADkaFTbjr0JGBpw"
Uri-Path: "tmid=567"
Content-Format: "application/dots+cbor"
```

```
{
  "ietf-dots-telemetry:telemetry": {
    "pre-or-ongoing-mitigation": [
      {
        "target": {
          "target-prefix": [
            "2001:db8::/32"
          ]
        }
      }
    ]
  }
}
```

Figure 39: PUT to Request Pre-or-Ongoing-Mitigation Telemetry

DOTS clients of the same domain can request to receive pre-or-ongoing-mitigation telemetry bound to the same target.

The DOTS client conveys the Observe Option set to '0' in the GET request to receive asynchronous notifications carrying pre-or-ongoing-mitigation telemetry data from the DOTS server. The GET request specifies a 'tmid' (Figure 40) or not (Figure 41).

```
Header: GET (Code=0.01)
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "tm"
Uri-Path: "cuid=dz6pHjaADkaFTbjr0JGBpw"
Uri-Path: "tmid=567"
Observe: 0
```

Figure 40: GET to Subscribe to Telemetry Asynchronous Notifications for a Specific 'tmid'

```
Header: GET (Code=0.01)
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "tm"
Uri-Path: "cuid=dz6pHjaADkaFTbjr0JGBpw"
Observe: 0
```

Figure 41: GET to Subscribe to Telemetry Asynchronous Notifications for All 'tmids'

The DOTS client can filter out the asynchronous notifications from the DOTS server by indicating one or more Uri-Query options in its GET request. A Uri-Query option can include the following parameters: 'target-prefix', 'target-port', 'target-protocol', 'target-fqdn', 'target-uri', 'alias-name', 'mid', and 'c' (content) (Section 4.2.4). Furthermore:

If more than one Uri-Query option is included in a request, these options are interpreted in the same way as when multiple target clauses are included in a message body.

If multiple values of a query parameter are to be included in a request, these values MUST be included in the same Uri-Query option and separated by a "," character without any spaces.

Range values (i.e., contiguous inclusive block) can be included for 'target-port', 'target-protocol', and 'mid' parameters by indicating two bound values separated by a "-" character.

Wildcard names (i.e., a name with the leftmost label is the "*" character) can be included in 'target-fqdn' or 'target-uri' parameters. DOTS clients MUST NOT include a name in which the "*" character is included in a label other than the leftmost label. "*.example.com" is an example of a valid wildcard name that can be included as a value of the 'target-fqdn' parameter in an Uri-Query option.

DOTS clients may also filter out the asynchronous notifications from

the DOTS server by indicating a specific source information. To that aim, a DOTS client may include 'source-prefix', 'source-port', or 'source-icmp-type' in a Uri-Query option. The same considerations (ranges, multiple values) specified for target clauses apply for source clauses. Special care SHOULD be taken when using these filters as some attacks may be hidden to the requesting DOTS client (e.g., the attack changes its source information).

Requests with invalid query types (e.g., not supported, malformed) by the DOTS server MUST be rejected by DOTS servers with a 4.00 (Bad Request).

An example of request to subscribe to asynchronous UDP telemetry notifications is shown in Figure 42. This filter will be applied for all 'tmids'.

```
Header: GET (Code=0.01)
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "tm"
Uri-Path: "cuid=dz6pHjaADkaFTbjr0JGBpw"
Uri-Query: "target-protocol=17"
Observe: 0
```

Figure 42: GET Request to Receive Telemetry Asynchronous Notifications Filtered using Uri-Query

The DOTS server will send asynchronous notifications to the DOTS client when an attack event is detected following similar considerations as in Section ~~5.4.2.1~~ **4.4.2.1** of

~~[I-D.boucaddair-dots-rfc8782-bis]~~ **[I-D.ietf-dots-rfc8782-bis]**.

An example of a ~~pre-or-ongoing-mitigation~~ **pre-or-ongoing-mitigation** telemetry notification is shown in Figure 43.

```
{
  "ietf-dots-telemetry:telemetry": {
    "pre-or-ongoing-mitigation": [
      {
        "tmid": 567,
        "target": {
          "target-prefix": [
            "2001:db8::1/128"
          ]
        },
        "target-protocol": [
          17
        ],
        "total-attack-traffic": [
          {
            "unit": "megabit-ps",
            "mid-percentile-g": "900"
          }
        ],
        "attack-detail": [
          {
            "vendor-id": 1234,
            "attack-id": 77,
            "start-time": "1957818434",
            "attack-severity": "high"
          }
        ]
      }
    ]
  }
}
```

Figure 43: Message Body of a Pre-or-Ongoing-Mitigation Telemetry Notification from the DOTS Server

A DOTS server sends the aggregate data for a target using 'total-attack-traffic' attribute. The aggregate assumes that Uri-Query filters are applied on the target. The DOTS server MAY include more granular data when needed (that is, 'total-attack-traffic-protocol' and 'total-attack-traffic-port'). If a port filter (or protocol filter) is included in a request, 'total-attack-traffic-protocol' (or 'total-attack-traffic-port') conveys the data with the port (or protocol) filter applied.

A DOTS server may aggregate pre-or-ongoing-mitigation data (e.g., 'top-talkers') for all targets of a domain, or when justified, send specific information (e.g., 'top-talkers') per individual targets.

The DOTS client may log pre-or-ongoing-mitigation telemetry data with an alert sent to an administrator or a network controller. The DOTS client may send a mitigation request if the attack cannot be handled locally.

A DOTS client that is not interested to receive pre-or-ongoing-mitigation telemetry data for a target MUST send a delete request similar to the one depicted in Figure 37.

8. DOTS Telemetry Mitigation Status Update

8.1. DOTS Clients to Servers Mitigation Efficacy DOTS Telemetry Attributes

The mitigation efficacy telemetry attributes can be signaled from DOTS clients to DOTS servers as part of the periodic mitigation efficacy updates to the server (Section 6.3.4 4.4.3 of ~~[I-D.boucadair-dots-rfc8782-bis]~~ **[I-D.ietf-dots-rfc8782-bis]**).

Total Attack Traffic: The overall attack traffic as observed from the DOTS client perspective during an active mitigation. See Figure 27.

Attack Details: The overall attack details as observed from the DOTS client perspective during an active mitigation. See Section 7.1.5.

The "ietf-dots-telemetry" YANG module (Section 10.1) augments the 'mitigation-scope' message type defined in "ietf-dots-signal" ~~[I-D.boucadair-dots-rfc8782-bis]~~ **[I-D.ietf-dots-rfc8782-bis]** so that these attributes can be signalled by a DOTS client in a mitigation efficacy update (Figure 44).

```
augment-structure /signal:dots-signal/signal:message-type
                  /signal:mitigation-scope/signal:scope: /dots-signal:dots-signal/dots-signal:message-type
                  /dots-signal:mitigation-scope/dots-signal:scope:

+-- total-attack-traffic* [unit]
|  +-- unit                unit
|  +-- low-percentile-g?   yang:gauge64
|  +-- mid-percentile-g?   yang:gauge64
|  +-- high-percentile-g?  yang:gauge64
|  +-- peak-g?            yang:gauge64
+-- attack-detail* [vendor-id attack-id]
   +-- vendor-id           uint32
   +-- attack-id           uint32
   +-- attack-description? string
   +-- attack-severity?    attack-severity
   +-- start-time?         uint64
   +-- end-time?           uint64
   +-- source-count
   |  +-- low-percentile-g? yang:gauge64
   |  +-- mid-percentile-g? yang:gauge64
   |  +-- high-percentile-g? yang:gauge64
   |  +-- peak-g?          yang:gauge64
   +-- top-talker
   |  +-- talker* [source-prefix]
   |  |  +-- spoofed-status?      boolean
   |  |  +-- source-prefix        inet:ip-prefix
   |  |  +-- source-port-range* [lower-port]
   |  |  |  +-- lower-port      inet:port-number
   |  |  |  +-- upper-port?     inet:port-number
   |  |  +-- source-icmp-type-range* [lower-type]
   |  |  |  +-- lower-type      uint8
   |  |  |  +-- upper-type?     uint8
   |  |  +-- total-attack-traffic* [unit]
   |  |  |  +-- unit            unit
   |  |  |  +-- low-percentile-g? yang:gauge64
   |  |  |  +-- mid-percentile-g? yang:gauge64
   |  |  |  +-- high-percentile-g? yang:gauge64
   |  |  |  +-- peak-g?        yang:gauge64
   |  +-- total-attack-connection
   |  |  +-- low-percentile-c
   |  |  |  +-- connection?      yang:gauge64
   |  |  |  +-- embryonic?       yang:gauge64
   |  |  |  +-- connection-ps?   yang:gauge64
   |  |  |  +-- request-ps?      yang:gauge64
   |  |  |  +-- partial-request-ps? yang:gauge64
   |  |  +-- mid-percentile-c
   |  |  |  ...
   |  |  +-- high-percentile-c
   |  |  |  ...
   |  |  +-- peak-c
   |  |  |  ...
```

Figure 44: Telemetry Efficacy Update Tree Structure

In order to signal telemetry data in a mitigation efficacy update, it is RECOMMENDED that the DOTS client has already established a DOTS telemetry setup session with the server in 'idle' time.

An example of an efficacy update with telemetry attributes is depicted in Figure 45.

```

Header: PUT (Code=0.03)
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "mitigate"
Uri-Path: "cuid=dz6pHjaADkaFTbjr0JGBpw"
Uri-Path: "mid=123"
If-Match:
Content-Format: "application/dots+cbor"

{
  "ietf-dots-signal-channel:mitigation-scope": {
    "scope": [
      {
        "alias-name": [
          "https1",
          "https2"
        ],
        "attack-status": "under-attack",
        "ietf-dots-telemetry:total-attack-traffic": [
          {
            "unit": "megabit-ps",
            "mid-percentile-g": "900"
          }
        ]
      }
    ]
  }
}

```

Figure 45: An Example of Mitigation Efficacy Update with Telemetry Attributes

8.2. DOTS Servers to Clients Mitigation Status DOTS Telemetry Attributes

The mitigation status telemetry attributes can be signaled from the DOTS server to the DOTS client as part of the periodic mitigation status update (Section ~~6.3.3~~ **4.4.2.2** of ~~[I-D.boucadair-dots-rfc8782-bis]~~ **[I-D.ietf-dots-rfc8782-bis]**). In particular, DOTS clients can receive asynchronous notifications of the attack details from DOTS servers using the Observe option defined in [RFC7641].

In order to make use of this feature, DOTS clients MUST establish a telemetry setup session with the DOTS server in 'idle' time and MUST set the 'server-originated-telemetry' attribute to 'true'.

DOTS servers MUST NOT include telemetry attributes in mitigation status updates sent to DOTS clients for which 'server-originated-telemetry' attribute is set to 'false'.

As defined in [RFC8612], the actual mitigation activities can include several countermeasure mechanisms. The DOTS server signals the current operational status of relevant countermeasures. A list of attacks detected by each countermeasure MAY also be included. The same attributes defined in Section 7.1.5 are applicable for describing the attacks detected and mitigated at the DOTS server domain.

The "ietf-dots-telemetry" YANG module (Section 10.1) augments the 'mitigation-scope' message type defined in "ietf-dots-signal" ~~[I-D.boucadair-dots-rfc8782-bis]~~ **[I-D.ietf-dots-rfc8782-bis]** with telemetry data as depicted in the following tree structure:

```

augment-structure /signal:dots-signal/signal:message-type
                  /signal:mitigation-scope/signal:scope: /dots-signal:dots-signal/dots-signal:message-type
                  /dots-signal:mitigation-scope/dots-signal:scope:
+-- (direction)?
|  +--:(server-to-client-only)
|  |  +-- total-traffic* [unit]
|  |  |  +-- unit
|  |  |  +-- low-percentile-g?   yang:gauge64
|  |  |  +-- mid-percentile-g?  yang:gauge64
|  |  |  +-- high-percentile-g? yang:gauge64
|  |  |  +-- peak-g?           yang:gauge64
|  |  +-- total-attack-connection
|  |  |  +-- low-percentile-c
|  |  |  |  +-- connection?      yang:gauge64
|  |  |  |  +-- embryonic?       yang:gauge64
|  |  |  |  +-- connection-ps?   yang:gauge64
|  |  |  |  +-- request-ps?      yang:gauge64
|  |  |  |  +-- partial-request-ps? yang:gauge64
|  |  |  +-- mid-percentile-c
|  |  |  |  ...
|  |  |  +-- high-percentile-c
|  |  |  |  ...
|  |  |  +-- peak-c

```

```

|
| ...
+-- total-attack-traffic* [unit]
| +-- unit unit
| +-- low-percentile-g? yang:gauge64
| +-- mid-percentile-g? yang:gauge64
| +-- high-percentile-g? yang:gauge64
| +-- peak-g? yang:gauge64
+-- attack-detail* [vendor-id attack-id]
+-- vendor-id uint32
+-- attack-id uint32
+-- attack-description? string
+-- attack-severity? attack-severity
+-- start-time? uint64
+-- end-time? uint64
+-- source-count
| +-- low-percentile-g? yang:gauge64
| +-- mid-percentile-g? yang:gauge64
| +-- high-percentile-g? yang:gauge64
| +-- peak-g? yang:gauge64
+-- top-talker
+-- talker* [source-prefix]
+-- spoofed-status? boolean
+-- source-prefix inet:ip-prefix
+-- source-port-range* [lower-port]
| +-- lower-port inet:port-number
| +-- upper-port? inet:port-number
+-- source-icmp-type-range* [lower-type]
| +-- lower-type uint8
| +-- upper-type? uint8
+-- total-attack-traffic* [unit]
| +-- unit unit
| +-- low-percentile-g? yang:gauge64
| +-- mid-percentile-g? yang:gauge64
| +-- high-percentile-g? yang:gauge64
| +-- peak-g? yang:gauge64
+-- total-attack-connection
+-- low-percentile-c
| +-- connection? yang:gauge64
| +-- embryonic? yang:gauge64
| +-- connection-ps? yang:gauge64
| +-- request-ps? yang:gauge64
| +-- partial-request-ps? yang:gauge64
+-- mid-percentile-c
| ...
+-- high-percentile-c
| ...
+-- peak-c
| ...

```

Figure 46 shows an example of an asynchronous notification of attack mitigation status from the DOTS server. This notification signals both the mid-percentile value of processed attack traffic and the peak percentile value of unique sources involved in the attack.

```

{
  "ietf-dots-signal-channel:mitigation-scope": {
    "scope": [
      {
        "mid": 12332,
        "mitigation-start": "1507818434",
        "alias-name": [
          "https1",
          "https2"
        ],
        "lifetime": 1600,
        "status": "attack-successfully-mitigated",
        "bytes-dropped": "134334555",
        "bps-dropped": "43344",
        "pkts-dropped": "333334444",
        "pps-dropped": "432432",
        "ietf-dots-telemetry:total-attack-traffic": [
          {
            "unit": "megabit-ps",
            "mid-percentile-g": "900"
          }
        ],
        "ietf-dots-telemetry:attack-detail": [
          {
            "vendor-id": 1234,
            "attack-id": 77,
            "source-count": {
              "peak-g": "10000"
            }
          }
        ]
      }
    ]
  }
}

```



```
}

```

Figure 46: Response Body of a Mitigation Status With Telemetry Attributes

DOTS clients can filter out the asynchronous notifications from the DOTS server by indicating one or more Uri-Query options in its GET request. A Uri-Query option can include the following parameters: 'target-prefix', 'target-port', 'target-protocol', 'target-fqdn', 'target-uri', 'alias-name', and 'c' (content) (Section 4.2.4). The considerations discussed in Section 7.3 MUST be followed to include multiple query values, ranges ('target-port', 'target-protocol'), and wildcard name ('target-fqdn', 'target-uri').

An example of request to subscribe to asynchronous notifications bound to the "http1" alias is shown in Figure 47.

```
Header: GET (Code=0.01)
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "mitigate"
Uri-Path: "cuid=dz6pHjaADkaFTbjr0JGBpw"
Uri-Path: "mid=12332"
Uri-Query: "target-alias=https1"
Observe: 0

```

Figure 47: GET Request to Receive Asynchronous Notifications Filtered using Uri-Query

If the target query does not match the target of the enclosed 'mid' as maintained by the DOTS server, the latter MUST respond with a 4.04 (Not Found) error Response Code. The DOTS server MUST NOT add a new observe entry if this query overlaps with an existing one.

9. Error Handling

A list of common CoAP errors that are implemented by DOTS servers are provided in Section 6.9 of ~~[I-D.boucadiar-dots-rfc8782-bis]~~ [I-D.ietf-dots-rfc8782-bis]. The following additional error cases apply for the telemetry extension:

- o 4.00 (Bad Request) is returned by the DOTS server when the DOTS client has sent a request that violates the DOTS telemetry extension.
- o 4.04 (Not Found) is returned by the DOTS server when the DOTS client is requesting a 'tsid' or 'tmid' that is not valid.
- o 4.00 (Bad Request) is returned by the DOTS server when the DOTS client has sent a request with invalid query types (e.g., not supported, malformed).
- o 4.04 (Not Found) is returned by the DOTS server when the DOTS client has sent a request with a target query that does not match the target of the enclosed 'mid' as maintained by the DOTS server.

10. YANG Modules

10.1. DOTS Signal Channel Telemetry YANG Module

This module uses types defined in [RFC6991] and [RFC8345].

```
<CODE BEGINS> file "ietf-dots-telemetry@2020-07-03.yang"
module ietf-dots-telemetry {
  yang-version 1.1;
  namespace "urn:ietf:params:xml:ns:yang:ietf-dots-telemetry";
  prefix dots-telemetry;

  import ietf-dots-signal-channel {
    prefix signal, dots-signal;
    reference
      "RFC UUUU: Distributed Denial-of-Service Open Threat Signaling
       (DOTS) Signal Channel Specification";
  }
  import ietf-dots-data-channel {
    prefix ietf-data, data-channel;
    reference
      "RFC 8783: Distributed Denial-of-Service Open Threat
       Signaling (DOTS) Data Channel Specification";
  }
  import ietf-yang-types {
    prefix yang;
    reference
      "Section 3 of RFC 6991";
  }
  import ietf-inet-types {
    prefix inet;
    reference
      "Section 4 of RFC 6991";
  }

```

```

}
import ietf-network-topology {
  prefix nt;
  reference
    "Section 6.2 of RFC 8345: A YANG Data Model for Network
    Topologies";
}
import ietf-yang-structure-ext {
  prefix sx;
  reference
    "RFC 8791: YANG Data Structure Extensions";
}

organization
  "IETF DDoS Open Threat Signaling (DOTS) Working Group";
contact
  "WG Web:  <https://datatracker.ietf.org/wg/dots/>
  WG List:  <mailto:dots@ietf.org>

  Author:   Mohamed Boucadair
            <mailto:mohamed.boucadair@orange.com>

  Author:   Konda, Tirumaleswar Reddy
            <mailto:TirumaleswarReddy_Konda@McAfee.com>";
description
  "This module contains YANG definitions for the signaling
  of DOTS telemetry exchanged between a DOTS client and
  a DOTS server by means of the DOTS signal channel.

  Copyright (c) 2020 IETF Trust and the persons identified as
  authors of the code. All rights reserved.

  Redistribution and use in source and binary forms, with or
  without modification, is permitted pursuant to, and subject
  to the license terms contained in, the Simplified BSD License
  set forth in Section 4.c of the IETF Trust's Legal Provisions
  Relating to IETF Documents
  (http://trustee.ietf.org/license-info).

  This version of this YANG module is part of RFC XXXX; see
  the RFC itself for full legal notices.";

revision 2020-07-03 {
  description
    "Initial revision.";
  reference
    "RFC XXXX: Distributed Denial-of-Service Open Threat
    Signaling (DOTS) Telemetry";
}

typedef attack-severity {
  type enumeration {
    enum none {
      value 1;
      description
        "No effect on the DOTS client domain.";
    }
    enum low {
      value 2;
      description
        "Minimal effect on the DOTS client domain.";
    }
    enum medium {
      value 3;
      description
        "A subset of DOTS client domain resources are
        out of service.";
    }
    enum high {
      value 4;
      description
        "The DOTS client domain is under extremely severe
        conditions.";
    }
    enum unknown {
      value 5;
      description
        "The impact of the attack is not known.";
    }
  }
  description
    "Enumeration for attack severity.";
  reference
    "RFC 7970: The Incident Object Description Exchange
    Format Version 2";
}

typedef unit-type {

```

```
type enumeration {
    enum packet-ps {
        value 1;
        description
            "Packets per second (pps).";
    }
    enum bit-ps {
        value 2;
        description
            "Bits per Second (bit/s).";
    }
    enum byte-ps {
        value 3;
        description
            "Bytes per second (Byte/s).";
    }
}
description
    "Enumeration to indicate which unit type is used.";
}
```

```
typedef unit {
    type enumeration {
        enum packet-ps {
            value 1;
            description
                "Packets per second (pps).";
        }
        enum bit-ps {
            value 2;
            description
                "Bits per Second (bps).";
        }
        enum byte-ps {
            value 3;
            description
                "Bytes per second (Bps).";
        }
        enum kilopacket-ps {
            value 4;
            description
                "Kilo packets per second (kpps).";
        }
        enum kilobit-ps {
            value 5;
            description
                "Kilobits per second (kbps).";
        }
        enum kilobyte-ps {
            value 6;
            description
                "Kilobytes per second (kBps).";
        }
        enum megapacket-ps {
            value 7;
            description
                "Mega packets per second (Mpps).";
        }
        enum megabit-ps {
            value 8;
            description
                "Megabits per second (Mbps).";
        }
        enum megabyte-ps {
            value 9;
            description
                "Megabytes per second (MBps).";
        }
        enum gigapacket-ps {
            value 10;
            description
                "Giga packets per second (Gpps).";
        }
        enum gigabit-ps {
            value 11;
            description
                "Gigabits per second (Gbps).";
        }
        enum gigabyte-ps {
            value 12;
            description
                "Gigabytes per second (GBps).";
        }
        enum terapacket-ps {
            value 13;
            description
                "Tera packets per second (Tpps).";
        }
    }
}
```

```
enum terabit-ps {
    value 14;
    description
        "Terabits per second (Tbps).";
}
enum terabyte-ps {
    value 15;
    description
        "Terabytes per second (TBps).";
}
}
description
    "Enumeration to indicate which unit is used.";
}

typedef interval {
    type enumeration {
        enum hour {
            value 1;
            description
                "Hour.";
        }
        enum day {
            value 2;
            description
                "Day.";
        }
        enum week {
            value 3;
            description
                "Week.";
        }
        enum month {
            value 4;
            description
                "Month.";
        }
    }
}
description
    "Enumeration to indicate the overall measurement period.";
}

typedef sample {
    type enumeration {
        enum second {
            value 1;
            description
                "A one second measurement period.";
        }
        enum 5-seconds {
            value 2;
            description
                "5 seconds measurement period.";
        }
        enum 30-seconds {
            value 3;
            description
                "30 seconds measurement period.";
        }
        enum minute {
            value 4;
            description
                "One minute measurement period.";
        }
        enum 5-minutes {
            value 5;
            description
                "5 minutes measurement period.";
        }
        enum 10-minutes {
            value 6;
            description
                "10 minutes measurement period.";
        }
        enum 30-minutes {
            value 7;
            description
                "30 minutes measurement period.";
        }
        enum hour {
            value 8;
            description
                "One hour measurement period.";
        }
    }
}
description
    "Enumeration to indicate the measurement period.";
}
```

```

typedef percentile {
    type decimal64 {
        fraction-digits 2;
    }
    description
        "The nth percentile of a set of data is the
         value at which n percent of the data is below it.";
}

typedef query-type {
    type enumeration {
        enum target-prefix {
            value 1;
            description
                "Query based on target prefix.";
        }
        enum target-port {
            value 2;
            description
                "Query based on target port number.";
        }
        enum target-protocol {
            value 3;
            description
                "Query based on target protocol.";
        }
        enum target-fqdn {
            value 4;
            description
                "Query based on target FQDN.";
        }
        enum target-uri {
            value 5;
            description
                "Query based on target URI.";
        }
        enum target-alias {
            value 6;
            description
                "Query based on target alias.";
        }
        enum mid {
            value 7;
            description
                "Query based on mitigation identifier (mid).";
        }
        enum source-prefix {
            value 8;
            description
                "Query based on source prefix.";
        }
        enum source-port {
            value 9;
            description
                "Query based on source port number.";
        }
        enum source-icmp-type {
            value 10;
            description
                "Query based on ICMP type";
        }
        enum content {
            value 11;
            description
                "Query based on 'c' Uri-Query option that is used
                 to control the selection of configuration
                 and non-configuration data nodes.";
            reference
                "Section 4.4.2 of RFC 8782.";
        }
    }
    description
        "Enumeration support for query types that can be used
         in a GET request to filter out data.";
}

grouping percentile-config {
    description
        "Configuration of low, mid, and high percentile values.";
    leaf measurement-interval {
        type interval;
        description
            "Defines the period on which percentiles are computed.";
    }
    leaf measurement-sample {
        type sample;
    }
}

```

```

    description
      "Defines the time distribution for measuring
       values that are used to compute percentiles.";
  }
  leaf low-percentile {
    type percentile;
    default "10.00";
    description
      "Low percentile. If set to '0', this means low-percentiles
       are disabled.";
  }
  leaf mid-percentile {
    type percentile;
    must '. >= ../low-percentile' {
      error-message
        "The mid-percentile must be greater than
         or equal to the low-percentile.";
    }
    default "50.00";
    description
      "Mid percentile. If set to the same value as low-percentiles,
       this means mid-percentiles are disabled.";
  }
  leaf high-percentile {
    type percentile;
    must '. >= ../mid-percentile' {
      error-message
        "The high-percentile must be greater than
         or equal to the mid-percentile.";
    }
    default "90.00";
    description
      "High percentile. If set to the same value as mid-percentiles,
       this means high-percentiles are disabled.";
  }
}

grouping percentile {
  description
    "Generic grouping for percentile.";
  leaf low-percentile-g {
    type yang:gauge64;
    description
      "Low percentile value.";
  }
  leaf mid-percentile-g {
    type yang:gauge64;
    description
      "Mid percentile value.";
  }
  leaf high-percentile-g {
    type yang:gauge64;
    description
      "High percentile value.";
  }
  leaf peak-g {
    type yang:gauge64;
    description
      "Peak value.";
  }
}

grouping unit-config {
  description
    "Generic grouping for unit configuration.";
  list unit-config {
    key "unit";
    description
      "Controls which unit types are allowed when sharing
       telemetry data.";
    leaf unit {
      type unit-type;
      description
        "Can be packet-ps, bit-ps, or byte-ps.";
    }
    leaf unit-status {
      type boolean;
      mandatory true;
      description
        "Enable/disable the use of the measurement unit type.";
    }
  }
}

grouping traffic-unit {
  description
    "Grouping of traffic as a function of the measurement unit.";
  leaf unit {

```

```

    type unit;
    description
        "The traffic can be measured using unit types: packet-ps,
        bit-ps, or byte-ps. DOTS agents auto-scale to the appropriate
        units (e.g., megabit-ps, kilobit-ps).";
    }
    uses percentile;
}

grouping traffic-unit-protocol {
    description
        "Grouping of traffic of a given transport protocol as
        a function of the measurement unit.";
    leaf protocol {
        type uint8;
        description
            "The transport protocol.
            Values are taken from the IANA Protocol Numbers registry:
            <https://www.iana.org/assignments/protocol-numbers/>.

            For example, this parameter contains 6 for TCP,
            17 for UDP, 33 for DCCP, or 132 for SCTP.";
    }
    uses traffic-unit;
}

grouping traffic-unit-port {
    description
        "Grouping of traffic bound to a port number as
        a function of the measurement unit.";
    leaf port {
        type inet:port-number;
        description
            "Port number.";
    }
    uses traffic-unit;
}

grouping total-connection-capacity {
    description
        "Total Connections Capacity. These data nodes are
        useful to detect resource consuming DDoS attacks";
    leaf connection {
        type uint64;
        description
            "The maximum number of simultaneous connections that
            are allowed to the target server.";
    }
    leaf connection-client {
        type uint64;
        description
            "The maximum number of simultaneous connections that
            are allowed to the target server per client.";
    }
    leaf embryonic {
        type uint64;
        description
            "The maximum number of simultaneous embryonic connections
            that are allowed to the target server. The term 'embryonic
            connection' refers to a connection whose connection handshake
            is not finished. Embryonic connection is only possible in
            connection-oriented transport protocols like TCP or SCTP.";
    }
    leaf embryonic-client {
        type uint64;
        description
            "The maximum number of simultaneous embryonic connections
            that are allowed to the target server per client.";
    }
    leaf connection-ps {
        type uint64;
        description
            "The maximum number of connections allowed per second
            to the target server.";
    }
    leaf connection-client-ps {
        type uint64;
        description
            "The maximum number of connections allowed per second
            to the target server per client.";
    }
    leaf request-ps {
        type uint64;
        description
            "The maximum number of requests allowed per second
            to the target server.";
    }
    leaf request-client-ps {

```

```

    type uint64;
    description
      "The maximum number of requests allowed per second
       to the target server per client.";
  }
  leaf partial-request-ps {
    type uint64;
    description
      "The maximum number of partial requests allowed per
       second to the target server.";
  }
  leaf partial-request-client-ps {
    type uint64;
    description
      "The maximum number of partial requests allowed per
       second to the target server per client.";
  }
}

grouping total-connection-capacity-protocol {
  description
    "Total Connections Capacity per protocol. These data nodes are
     useful to detect resource consuming DDoS attacks.";

  leaf protocol {
    type uint8;
    description
      "The transport protocol.
       Values are taken from the IANA Protocol Numbers registry:
       <https://www.iana.org/assignments/protocol-numbers/>.";
  }
  uses total-connection-capacity;
}

grouping connection {
  description
    "A set of data nodes which represent the attack
     characteristics";
  leaf connection {
    type yang:gauge64;
    description
      "The number of simultaneous attack connections to
       the target server.";
  }
  leaf embryonic {
    type yang:gauge64;
    description
      "The number of simultaneous embryonic connections to
       the target server.";
  }
  leaf connection-ps {
    type yang:gauge64;
    description
      "The number of attack connections per second to
       the target server.";
  }
  leaf request-ps {
    type yang:gauge64;
    description
      "The number of attack requests per second to
       the target server.";
  }
  leaf partial-request-ps {
    type yang:gauge64;
    description
      "The number of attack partial requests to
       the target server.";
  }
}

grouping connection-percentile {
  description
    "Total attack connections.";
  container low-percentile-c {
    description
      "Low percentile of attack connections.";
    uses connection;
  }
  container mid-percentile-c {
    description
      "Mid percentile of attack connections.";
    uses connection;
  }
  container high-percentile-c {
    description
      "High percentile of attack connections.";
    uses connection;
  }
}

```



```
    container peak-c {
      description
        "Peak attack connections.";
      uses connection;
    }
  }

  grouping connection-protocol {
    description
      "Total attack connections.";
    leaf protocol {
      type uint8;
      description
        "The transport protocol.
        Values are taken from the IANA Protocol Numbers registry:
        <https://www.iana.org/assignments/protocol-numbers/>.";
    }
    uses connection;
  }

  grouping connection-port {
    description
      "Total attack connections per port number.";
    leaf port {
      type inet:port-number;
      description
        "Port number.";
    }
    uses connection-protocol;
  }

  grouping connection-protocol-percentile {
    description
      "Total attack connections per protocol.";
    list low-percentile-1 {
      key "protocol";
      description
        "Low percentile of attack connections per protocol.";
      uses connection-protocol;
    }
    list mid-percentile-1 {
      key "protocol";
      description
        "Mid percentile of attack connections per protocol.";
      uses connection-protocol;
    }
    list high-percentile-1 {
      key "protocol";
      description
        "High percentile of attack connections per protocol.";
      uses connection-protocol;
    }
    list peak-1 {
      key "protocol";
      description
        "Peak attack connections per protocol.";
      uses connection-protocol;
    }
  }

  grouping connection-protocol-port-percentile {
    description
      "Total attack connections per port number.";
    list low-percentile-1 {
      key "protocol port";
      description
        "Low percentile of attack connections per port number.";
      uses connection-port;
    }
    list mid-percentile-1 {
      key "protocol port";
      description
        "Mid percentile of attack connections per port number.";
      uses connection-port;
    }
    list high-percentile-1 {
      key "protocol port";
      description
        "High percentile of attack connections per port number.";
      uses connection-port;
    }
    list peak-1 {
      key "protocol port";
      description
        "Peak attack connections per port number.";
      uses connection-port;
    }
  }
```

```

}

grouping attack-detail {
  description
    "Various details that describe the on-going
    attacks that need to be mitigated by the DOTS server.
    The attack details need to cover well-known and common attacks
    (such as a SYN Flood) along with new emerging or vendor-specific
    attacks.";
  leaf vendor-id {
    type uint32;
    description
      "Vendor ID is a security vendor's Enterprise Number.";
  }
  leaf attack-id {
    type uint32;
    description
      "Unique identifier assigned by the vendor for the attack.";
  }
  leaf attack-description {
    type string;
    description
      "Textual representation of attack description. Natural Language
      Processing techniques (e.g., word embedding) can possibly be
      used to map the attack description to an attack type.";
  }
  leaf attack-severity {
    type attack-severity;
    description
      "Severity level of an attack. How this level is determined
      is implementation-specific.";
  }
  leaf start-time {
    type uint64;
    description
      "The time the attack started. Start time is represented in
      seconds relative to 1970-01-01T00:00:00Z in UTC time.";
  }
  leaf end-time {
    type uint64;
    description
      "The time the attack ended. End time is represented in seconds
      relative to 1970-01-01T00:00:00Z in UTC time.";
  }
  container source-count {
    description
      "Indicates the count of unique sources involved
      in the attack.";
    uses percentile;
  }
}

grouping top-talker-aggregate {
  description
    "Top attack sources.";
  list talker {
    key "source-prefix";
    description
      "IPv4 or IPv6 prefix identifying the attacker(s).";
    leaf spoofed-status {
      type boolean;
      description
        "Indicates whether this address is spoofed.";
    }
    leaf source-prefix {
      type inet:ip-prefix;
      description
        "IPv4 or IPv6 prefix identifying the attacker(s).";
    }
  }
  list source-port-range {
    key "lower-port";
    description
      "Port range. When only lower-port is
      present, it represents a single port number.";
    leaf lower-port {
      type inet:port-number;
      mandatory true;
      description
        "Lower port number of the port range.";
    }
    leaf upper-port {
      type inet:port-number;
      must '. >= ../lower-port' {
        error-message
          "The upper port number must be greater than
          or equal to lower port number.";
      }
    }
    description

```

```

        "Upper port number of the port range.";
    }
}
list source-icmp-type-range {
    key "lower-type";
    description
        "ICMP type range. When only lower-type is
        present, it represents a single ICMP type.";
    leaf lower-type {
        type uint8;
        mandatory true;
        description
            "Lower ICMP type of the ICMP type range.";
    }
    leaf upper-type {
        type uint8;
        must '. >= ../lower-type' {
            error-message
                "The upper ICMP type must be greater than
                or equal to lower ICMP type.";
        }
        description
            "Upper type of the ICMP type range.";
    }
}
list total-attack-traffic {
    key "unit";
    description
        "Total attack traffic issued from this source.";
    uses traffic-unit;
}
container total-attack-connection {
    description
        "Total attack connections issued from this source.";
    uses connection-percentile;
}
}
}

grouping top-talker {
    description
        "Top attack sources.";
    list talker {
        key "source-prefix";
        description
            "IPv4 or IPv6 prefix identifying the attacker(s).";
        leaf spoofed-status {
            type boolean;
            description
                "Indicates whether this address is spoofed.";
        }
        leaf source-prefix {
            type inet:ip-prefix;
            description
                "IPv4 or IPv6 prefix identifying the attacker(s).";
        }
    }
    list source-port-range {
        key "lower-port";
        description
            "Port range. When only lower-port is
            present, it represents a single port number.";
        leaf lower-port {
            type inet:port-number;
            mandatory true;
            description
                "Lower port number of the port range.";
        }
        leaf upper-port {
            type inet:port-number;
            must '. >= ../lower-port' {
                error-message
                    "The upper port number must be greater than
                    or equal to lower port number.";
            }
            description
                "Upper port number of the port range.";
        }
    }
}
list source-icmp-type-range {
    key "lower-type";
    description
        "ICMP type range. When only lower-type is
        present, it represents a single ICMP type.";
    leaf lower-type {
        type uint8;
        mandatory true;
        description

```

```

        "Lower ICMP type of the ICMP type range.";
    }
    leaf upper-type {
        type uint8;
        must '. >= ../lower-type' {
            error-message
                "The upper ICMP type must be greater than
                 or equal to lower ICMP type.";
        }
        description
            "Upper type of the ICMP type range.";
    }
}
list total-attack-traffic {
    key "unit";
    description
        "Total attack traffic issued from this source.";
    uses traffic-unit;
}
container total-attack-connection {
    description
        "Total attack connections issued from this source.";
    uses connection-protocol-percentile;
}
}
}

```

```

grouping baseline {
    description
        "Grouping for the telemetry baseline.";
    uses ietf-data-target, data-channel:target;
    leaf-list alias-name {
        type string;
        description
            "An alias name that points to a resource.";
    }
    list total-traffic-normal {
        key "unit";
        description
            "Total traffic normal baselines.";
        uses traffic-unit;
    }
    list total-traffic-normal-per-protocol {
        key "unit protocol";
        description
            "Total traffic normal baselines per protocol.";
        uses traffic-unit-protocol;
    }
    list total-traffic-normal-per-port {
        key "unit port";
        description
            "Total traffic normal baselines per port number.";
        uses traffic-unit-port;
    }
    list total-connection-capacity {
        key "protocol";
        description
            "Total connection capacity.";

        uses total-connection-capacity-protocol;
    }
    list total-connection-capacity-per-port {
        key "protocol port";
        description
            "Total connection capacity per port number.";
        leaf port {
            type inet:port-number;
            description
                "The target port number.";
        }
        uses total-connection-capacity-protocol;
    }
}
}

```

```

grouping pre-or-ongoing-mitigation {
    description
        "Grouping for the telemetry data.";
    list total-traffic {
        key "unit";
        description
            "Total traffic.";
        uses traffic-unit;
    }
    list total-traffic-protocol {
        key "unit protocol";
        description
            "Total traffic per protocol.";
        uses traffic-unit-protocol;
    }
}

```

```

}
list total-traffic-port {
  key "unit port";
  description
    "Total traffic per port.";
  uses traffic-unit-port;
}
list total-attack-traffic {
  key "unit";
  description
    "Total attack traffic.";
  uses traffic-unit-protocol;
}
list total-attack-traffic-protocol {
  key "unit protocol";
  description
    "Total attack traffic per protocol.";
  uses traffic-unit-protocol;
}
list total-attack-traffic-port {
  key "unit port";
  description
    "Total attack traffic per port.";
  uses traffic-unit-port;
}
container total-attack-connection {
  description
    "Total attack connections.";
  uses connection-protocol-percentile;
}
container total-attack-connection-port {
  description
    "Total attack connections.";
  uses connection-protocol-port-percentile;
}
list attack-detail {
  key "vendor-id attack-id";
  description
    "Provides a set of attack details.";
  uses attack-detail;
  container top-talker {
    description
      "Lists the top attack sources.";
    uses top-talker;
  }
}
}

sx:augment-structure "/signal:dots-signal/signal:message-type/" "/dots-signal:dots-signal/dots-signal:message-type/"
  + "signal:mitigation-scope/signal:scope" "dots-signal:mitigation-scope/dots-signal:scope" {
  description
    "Extends mitigation scope with telemetry update data.";
  choice direction {
    description
      "Indicates the communication direction in which the
      data nodes can be included.";
    case server-to-client-only {
      description
        "These data nodes appear only in a mitigation message
        sent from the server to the client.";
      list total-traffic {
        key "unit";
        description
          "Total traffic.";
        uses traffic-unit;
      }
      container total-attack-connection {
        description
          "Total attack connections.";
        uses connection-percentile;
      }
    }
  }
  list total-attack-traffic {
    key "unit";
    description
      "Total attack traffic.";
    uses traffic-unit;
  }
  list attack-detail {
    key "vendor-id attack-id";
    description
      "Attack details";
    uses attack-detail;
    container top-talker {
      description
        "Top attack sources.";
      uses top-talker-aggregate;
    }
  }
}

```

```

    }
  }
}
sx:structure dots-telemetry {
  description
    "Main structure for DOTS telemetry messages.";
  choice telemetry-message-type {
    description
      "Can be a telemetry-setup or telemetry data.";
    case telemetry-setup {
      description
        "Indicates the message is about telemetry.";
      choice direction {
        description
          "Indicates the communication direction in which the
            data nodes can be included.";
        case server-to-client-only {
          description
            "These data nodes appear only in a mitigation message
              sent from the server to the client.";
          container max-config-values {
            description
              "Maximum acceptable configuration values.";
            uses percentile-config;
            leaf server-originated-telemetry {
              type boolean;
              description
                "Indicates whether the DOTS server can be instructed
                  to send pre-or-ongoing-mitigation telemetry. If set
                  to FALSE or the data node is not present, this is
                  an indication that the server does not support this
                  capability.";
            }
            leaf telemetry-notify-interval {
              type uint32 {
                range "1 .. 3600";
              }
              must ". >= ../min-config-values"
                + "/telemetry-notify-interval" {
                error-message
                  "The value must be greater than or equal
                    to the telemetry-notify-interval in the
                    min-config-values";
              }
              units "seconds";
              description
                "Minimum number of seconds between successive
                  telemetry notifications.";
            }
          }
        }
      }
    }
    container min-config-values {
      description
        "Minimum acceptable configuration values.";
      uses percentile-config;
      leaf telemetry-notify-interval {
        type uint32 {
          range "1 .. 3600";
        }
        units "seconds";
        description
          "Minimum number of seconds between successive
            telemetry notifications.";
      }
    }
  }
  container supported-units {
    description
      "Supported units and default activation status.";
    uses unit-config;
  }
  leaf-list query-type {
    type query-type;
    description
      "Indicates which query types are supported by
        the server.";
  }
}
}
list telemetry {
  description
    "The telemetry data per DOTS client.";
  choice direction {
    description
      "Indicates the communication direction in which the
        data nodes can be included.";
    case server-to-client-only {
      description
        "These data nodes appear only in a mitigation message

```

```

        sent from the server to the client.";
    leaf tsid {
        type uint32;
        description
            "An identifier for the DOTS telemetry setup
            data.";
    }
}
}
choice setup-type {
    description
        "Can be a mitigation configuration, a pipe capacity,
        or baseline message.";
    case telemetry-config {
        description
            "Uses to set low, mid, and high percentile values.";
        container current-config {
            description
                "Current configuration values.";
            uses percentile-config;
            uses unit-config;
            leaf server-originated-telemetry {
                type boolean;
                description
                    "Used by a DOTS client to enable/disable whether it
                    accepts pre-or-ongoing-mitigation telemetry from
                    the DOTS server.";
            }
            leaf telemetry-notify-interval {
                type uint32 {
                    range "1 .. 3600";
                }
                units "seconds";
                description
                    "Minimum number of seconds between successive
                    telemetry notifications.";
            }
        }
    }
}
}
case pipe {
    description
        "Total pipe capacity of a DOTS client domain";
    list total-pipe-capacity {
        key "link-id unit";
        description
            "Total pipe capacity of a DOTS client domain.";
        leaf link-id {
            type nt:link-id;
            description
                "Identifier of an interconnection link.";
        }
        leaf capacity {
            type uint64;
            mandatory true;
            description
                "Pipe capacity.";
        }
        leaf unit {
            type unit;
            description
                "The traffic can be measured using unit types:
                packets per second (PPS), Bits per Second (BPS),
                and/or bytes per second. DOTS agents auto-scale
                to the appropriate units (e.g., megabit-ps,
                kilobit-ps).";
        }
    }
}
}
case baseline {
    description
        "Traffic baseline information";
    list baseline {
        key "id";
        description
            "Traffic baseline information";
        leaf id {
            type uint32;
            must '. >= 1';
            description
                "A baseline entry identifier.";
        }
        uses baseline;
    }
}
}
}

```

```

    }
    case telemetry {
      description
        "Indicates the message is about telemetry.";
      list pre-or-ongoing-mitigation {
        description
          "Pre-or-ongoing-mitigation telemetry per DOTS client.";
        choice direction {
          description
            "Indicates the communication direction in which the
             data nodes can be included.";
          case server-to-client-only {
            description
              "These data nodes appear only in a mitigation message
               sent from the server to the client.";
            leaf tmid {
              type uint32;
              description
                "An identifier to uniquely demux telemetry data sent
                 using the same message.";
            }
          }
        }
      }
      container target {
        description
          "Indicates the target.";
        uses ietf-data:target data-channel:target;
        leaf-list alias-name {
          type string;
          description
            "An alias name that points to a resource.";
        }
        leaf-list mid-list {
          type uint32;
          description
            "Reference a list of associated mitigation requests.";
        }
      }
      uses pre-or-ongoing-mitigation;
    }
  }
}
}
}
<CODE ENDS>

```

10.2. Vendor Attack Mapping Details YANG Module

```

<CODE BEGINS> file "ietf-dots-mapping@2020-06-26.yang"
module ietf-dots-mapping {
  yang-version 1.1;
  namespace "urn:ietf:params:xml:ns:yang:ietf-dots-mapping";
  prefix dots-mapping;

  import ietf-dots-data-channel {
    prefix ietf-data data-channel;
    reference
      "RFC 8783: Distributed Denial-of-Service Open Threat
       Signaling (DOTS) Data Channel Specification";
  }

  organization
    "IETF DDoS Open Threat Signaling (DOTS) Working Group";
  contact
    "WG Web:  <https://datatracker.ietf.org/wg/dots/>
     WG List: <mailto:dots@ietf.org>

     Author: Mohamed Boucadair
              <mailto:mohamed.boucadair@orange.com>

     Author: Jon Shallow
              <mailto:supjps-ietf@jpshallow.com>";
  description
    "This module contains YANG definitions for the sharing
     DDoS attack mapping details between a DOTS client and
     a DOTS server, by means of the DOTS data channel.

     Copyright (c) 2020 IETF Trust and the persons identified as
     authors of the code. All rights reserved.

     Redistribution and use in source and binary forms, with or
     without modification, is permitted pursuant to, and subject
     to the license terms contained in, the Simplified BSD License
     set forth in Section 4.c of the IETF Trust's Legal Provisions
     Relating to IETF Documents
     (http://trustee.ietf.org/license-info).

     This version of this YANG module is part of RFC XXXX; see
     the RFC itself for full legal notices.";

```



```

revision 2020-06-26 {
  description
    "Initial revision.";
  reference
    "RFC XXXX: Distributed Denial-of-Service Open Threat
      Signaling (DOTS) Telemetry";
}

feature dots-telemetry {
  description
    "This feature indicates that DOTS telemetry data can be
      shared between DOTS clients and servers.";
}

grouping attack-mapping {
  description
    "A set of information used for sharing vendor attack mapping
      information with a peer.";
  list vendor {
    key "vendor-id";
    description
      "Vendor attack mapping information of the client/server";
    leaf vendor-id {
      type uint32;
      description
        "Vendor ID is a security vendor's Enterprise Number.";
    }
    leaf vendor-name {
      type string;
      description
        "The name of the vendor (e.g., company A).";
    }
    leaf last-updated {
      type uint64;
      mandatory true;
      description
        "The time the mapping table was updated. It is represented
          in seconds relative to 1970-01-01T00:00:00Z in UTC time.";
    }
    list attack-mapping {
      key "attack-id";
      description
        "Attack mapping details.";
      leaf attack-id {
        type uint32;
        description
          "Unique identifier assigned by the vendor for the attack.";
      }
      leaf attack-description {
        type string;
        mandatory true;
        description
          "Textual representation of attack description. Natural
            Language Processing techniques (e.g., word embedding)
            can possibly be used to map the attack description to
            an attack type.";
      }
    }
  }
}

augment "/ietf-data:dots-data/ietf-data:dots-client" "/data-channel:dots-data/data-channel:dots-client" {
  if-feature "dots-telemetry";
  description
    "Augments the data channel with a vendor attack
      mapping table of the DOTS client.";
  container vendor-mapping {
    description
      "Used by DOTS clients to share their vendor
        attack mapping information with DOTS servers.";
    uses attack-mapping;
  }
}

augment "/ietf-data:dots-data/ietf-data:capabilities" "/data-channel:dots-data/data-channel:capabilities" {
  if-feature "dots-telemetry";
  description
    "Augments the DOTS server capabilities with a
      parameter to indicate whether they can share
      attack mapping details.";
  leaf vendor-mapping-enabled {
    type boolean;
    config false;
    description
      "Indicates that the server supports sharing
        attack vendor mapping details with DOTS clients.";
  }
}

```

```

    }
  }
  augment "/ietf-data:dots-data" "/data-channel:dots-data" {
    if-feature "dots-telemetry";
    description
      "Augments the data channel with a vendor attack
       mapping table of the DOTS server.";
    container vendor-mapping {
      config false;
      description
        "Includes the list of vendor attack mapping details
         that will be shared upon request with DOTS clients.";
      uses attack-mapping;
    }
  }
}
<CODE ENDS>

```

11. YANG/JSON Mapping Parameters to CBOR

All DOTS telemetry parameters in the payload of the DOTS signal channel MUST be mapped to CBOR types as shown in the following table:

- o Implementers may use the values in: <https://github.com/boucadair/draft-dots-telemetry/blob/master/mapping-table.txt>

Parameter Name	YANG Type	CBOR Key	CBOR Major Type & Information	JSON Type
tsid	uint32	TBA1	0 unsigned	Number
telemetry	container	TBA2	5 map	Object
low-percentile	decimal64	TBA3	6 tag 4	String
			[-2, integer]	
mid-percentile	decimal64	TBA4	6 tag 4	String
			[-2, integer]	
high-percentile	decimal64	TBA5	6 tag 4	String
			[-2, integer]	
unit-config	list	TBA6	4 array	Array
unit	enumeration	TBA7	0 unsigned	String
unit-status	boolean	TBA8	7 bits 20	False
			7 bits 21	True
total-pipe-capability	list	TBA9	4 array	Array
link-id	string	TBA10	3 text string	String
pre-or-ongoing-mitigation	list	TBA11	4 array	Array
total-traffic-normal	list	TBA12	4 array	Array
low-percentile-g	yang:gauge64	TBA13	0 unsigned	String
mid-percentile-g	yang:gauge64	TBA14	0 unsigned	String
high-percentile-g	yang:gauge64	TBA15	0 unsigned	String
peak-g	yang:gauge64	TBA16	0 unsigned	String
total-attack-traffic	list	TBA17	4 array	Array
total-traffic	list	TBA18	4 array	Array
total-connection-capacity	list	TBA19	4 array	Array
connection	uint64	TBA20	0 unsigned	String
connection-client	uint64	TBA21	0 unsigned	String
embryonic	uint64	TBA22	0 unsigned	String
embryonic-client	uint64	TBA23	0 unsigned	String
connection-ps	uint64	TBA24	0 unsigned	String
connection-client-ps	uint64	TBA25	0 unsigned	String
request-ps	uint64	TBA26	0 unsigned	String
request-client-ps	uint64	TBA27	0 unsigned	String
partial-request-ps	uint64	TBA28	0 unsigned	String
partial-request-client-ps	uint64	TBA29	0 unsigned	String
total-attack-connection	container	TBA30	5 map	Object
low-percentile-l	list	TBA31	4 array	Array
mid-percentile-l	list	TBA32	4 array	Array
high-percentile-l	list	TBA33	4 array	Array
peak-l	list	TBA34	4 array	Array
attack-detail	list	TBA35	4 array	Array
id	uint32	TBA36	0 unsigned	Number
attack-id	uint32	TBA37	0 unsigned	Number
attack-description	string	TBA38	3 text string	String
attack-severity	enumeration	TBA39	0 unsigned	String
start-time	uint64	TBA40	0 unsigned	String
end-time	uint64	TBA41	0 unsigned	String
source-count	container	TBA42	5 map	Object
top-talker	container	TBA43	5 map	Object
spoofed-status	boolean	TBA44	7 bits 20	False
			7 bits 21	True
low-percentile-c	container	TBA45	5 map	Object
mid-percentile-c	container	TBA46	5 map	Object
high-percentile-c	container	TBA47	5 map	Object

peak-c	container	TBA48	5 map	Object
baseline	container	TBA49	5 map	Object
current-config	container	TBA50	5 map	Object
max-config-values	container	TBA51	5 map	Object
min-config-values	container	TBA52	5 map	Object
supported-units	container	TBA53	5 map	Object
server-originated-telemetry	boolean	TBA54	7 bits 20	False
telemetry-notify-interval	uint32	TBA55	7 bits 21	True
tmid	uint32	TBA56	0 unsigned	Number
measurement-interval	enumeration	TBA57	0 unsigned	String
measurement-sample	enumeration	TBA58	0 unsigned	String
talker	list	TBA59	4 array	Array
source-prefix	inet:	TBA60	3 text string	String
	ip-prefix			
mid-list	leaf-list	TBA61	4 array	Array
	uint32		0 unsigned	Number
source-port-range	list	TBA62	4 array	Array
source-icmp-type-range	list	TBA63	4 array	Array
lower-type	uint8	TBA64	0 unsigned	Number
upper-type	uint8	TBA65	0 unsigned	Number
target	container	TBA66	5 map	Object
capacity	uint64	TBA67	0 unsigned	String
protocol	uint8	TBA68	0 unsigned	Number
total-traffic-normal-per-protocol	list	TBA69	4 array	Array
total-traffic-normal-per-port	list	TBA70	4 array	Array
total-connection-capacity-per-port	list	TBA71	4 array	Array
total-traffic-protocol	list	TBA72	4 array	Array
total-traffic-port	list	TBA73	4 array	Array
total-attack-traffic-protocol	list	TBA74	4 array	Array
total-attack-traffic-port	list	TBA75	4 array	Array
total-attack-connection-port	list	TBA76	4 array	Array
port	inet:			
	port-number	TBA77	0 unsigned	Number
query-type	leaf-list	TBA78	4 array	Array
			0 unsigned	String
vendor-id	uint32	TBA79	0 unsigned	Number
ietf-dots-telemetry:telemetry-setup	container	TBA80	5 map	Object
ietf-dots-telemetry:total-traffic	list	TBA81	4 array	Array
ietf-dots-telemetry:total-attack-traffic	list	TBA82	4 array	Array
ietf-dots-telemetry:total-attack-connection	container	TBA83	5 map	Object
ietf-dots-telemetry:attack-detail	list	TBA84	4 array	Array

12. IANA Considerations

12.1. DOTS Signal Channel CBOR Key Values

This specification registers the DOTS telemetry attributes in the IANA "DOTS Signal Channel CBOR Key Values" registry [Key-Map].

The DOTS telemetry attributes defined in this specification are comprehension-optional parameters.

- o Note to the RFC Editor: CBOR keys are assigned from the 128-255 range.

Parameter Name	CBOR Key	CBOR Major Type	Change Controller	Specification Document(s)
tsid	TBA1	0	IESG	[RFCXXXX]
telemetry	TBA2	5	IESG	[RFCXXXX]
low-percentile	TBA3	6tag4	IESG	[RFCXXXX]
mid-percentile	TBA4	6tag4	IESG	[RFCXXXX]
high-percentile	TBA5	6tag4	IESG	[RFCXXXX]
unit-config	TBA6	4	IESG	[RFCXXXX]
unit	TBA7	0	IESG	[RFCXXXX]
unit-status	TBA8	7	IESG	[RFCXXXX]
total-pipe-capability	TBA9	4	IESG	[RFCXXXX]
link-id	TBA10	3	IESG	[RFCXXXX]
pre-or-ongoing-	TBA11	4	IESG	[RFCXXXX]

mitigation					
total-traffic-normal	TBA12	4	IESG	[RFCXXXX]	
low-percentile-g	TBA13	0	IESG	[RFCXXXX]	
mid-percentile-g	TBA14	0	IESG	[RFCXXXX]	
high-percentile-g	TBA15	0	IESG	[RFCXXXX]	
peak-g	TBA16	0	IESG	[RFCXXXX]	
total-attack-traffic	TBA17	4	IESG	[RFCXXXX]	
total-traffic	TBA18	4	IESG	[RFCXXXX]	
total-connection-	TBA19	4	IESG	[RFCXXXX]	
capacity					
connection	TBA20	0	IESG	[RFCXXXX]	
connection-client	TBA21	0	IESG	[RFCXXXX]	
embryonic	TBA22	0	IESG	[RFCXXXX]	
embryonic-client	TBA23	0	IESG	[RFCXXXX]	
connection-ps	TBA24	0	IESG	[RFCXXXX]	
connection-client-ps	TBA25	0	IESG	[RFCXXXX]	
request-ps	TBA26	0	IESG	[RFCXXXX]	
request-client-ps	TBA27	0	IESG	[RFCXXXX]	
partial-request-ps	TBA28	0	IESG	[RFCXXXX]	
partial-request-	TBA29	0	IESG	[RFCXXXX]	
client-ps					
total-attack-	TBA30	5	IESG	[RFCXXXX]	
connection					
low-percentile-l	TBA31	4	IESG	[RFCXXXX]	
mid-percentile-l	TBA32	4	IESG	[RFCXXXX]	
high-percentile-l	TBA33	4	IESG	[RFCXXXX]	
peak-l	TBA34	4	IESG	[RFCXXXX]	
attack-detail	TBA35	4	IESG	[RFCXXXX]	
id	TBA36	0	IESG	[RFCXXXX]	
attack-id	TBA37	0	IESG	[RFCXXXX]	
attack-description	TBA38	3	IESG	[RFCXXXX]	
attack-severity	TBA39	0	IESG	[RFCXXXX]	
start-time	TBA40	0	IESG	[RFCXXXX]	
end-time	TBA41	0	IESG	[RFCXXXX]	
source-count	TBA42	5	IESG	[RFCXXXX]	
top-talker	TBA43	5	IESG	[RFCXXXX]	
spoofed-status	TBA44	7	IESG	[RFCXXXX]	
low-percentile-c	TBA45	5	IESG	[RFCXXXX]	
mid-percentile-c	TBA46	5	IESG	[RFCXXXX]	
high-percentile-c	TBA47	5	IESG	[RFCXXXX]	
peak-c	TBA48	5	IESG	[RFCXXXX]	
ietf-dots-signal-cha	TBA49	5	IESG	[RFCXXXX]	
current-config	TBA50	5	IESG	[RFCXXXX]	
max-config-value	TBA51	5	IESG	[RFCXXXX]	
min-config-values	TBA52	5	IESG	[RFCXXXX]	
supported-units	TBA55	5	IESG	[RFCXXXX]	
server-originated-	TBA54	7	IESG	[RFCXXXX]	
telemetry					
telemetry-notify-	TBA55	0	IESG	[RFCXXXX]	
interval					
tmid	TBA56	0	IESG	[RFCXXXX]	
measurement-interval	TBA57	0	IESG	[RFCXXXX]	
measurement-sample	TBA58	0	IESG	[RFCXXXX]	
talker	TBA59	0	IESG	[RFCXXXX]	
source-prefix	TBA60	0	IESG	[RFCXXXX]	
mid-list	TBA61	4	IESG	[RFCXXXX]	
source-port-range	TBA62	4	IESG	[RFCXXXX]	
source-icmp-type-	TBA63	4	IESG	[RFCXXXX]	
range					
lower-type	TBA64	0	IESG	[RFCXXXX]	
upper-type	TBA65	0	IESG	[RFCXXXX]	
target	TBA66	5	IESG	[RFCXXXX]	
capacity	TBA67	0	IESG	[RFCXXXX]	
protocol	TBA68	0	IESG	[RFCXXXX]	
total-traffic-	TBA69	4	IESG	[RFCXXXX]	
normal-per-protocol					
total-traffic-	TBA70	4	IESG	[RFCXXXX]	
normal-per-port					
total-connection-	TBA71	4	IESG	[RFCXXXX]	
capacity-per-port					
total-traffic-	TBA72	4	IESG	[RFCXXXX]	
-protocol					
total-traffic-port	TBA73	4	IESG	[RFCXXXX]	
total-attack-	TBA74	4	IESG	[RFCXXXX]	
traffic-protocol					
total-attack-	TBA75	4	IESG	[RFCXXXX]	
traffic-port					
total-attack-	TBA76	4	IESG	[RFCXXXX]	
connection-port					
port	TBA77	0	IESG	[RFCXXXX]	
query-type	TBA78	4	IESG	[RFCXXXX]	
vendor-id	TBA79	0	IESG	[RFCXXXX]	
ietf-dots-telemetry:	TBA80	5	IESG	[RFCXXXX]	
telemetry-setup					
ietf-dots-telemetry:	TBA81	0	IESG	[RFCXXXX]	
total-traffic					
ietf-dots-telemetry:	TBA82	0	IESG	[RFCXXXX]	
total-attack-traffic					

ietf-dots-telemetry:	TBA83	0	IESG	[RFCXXXX]	
total-attack-					
connection					
ietf-dots-telemetry:	TBA84	4	IESG	[RFCXXXX]	
attack-detail					
+-----+-----+-----+-----+-----+					

12.2. DOTS Signal Channel Conflict Cause Codes

This specification requests IANA to assign a new code from the "DOTS Signal Channel Conflict Cause Codes" registry [Cause].

+-----+	+-----+	+-----+	+-----+	+-----+
Code	Label	Description	Reference	
+-----+-----+-----+-----+-----+				
TBA	overlapping-pipes	Overlapping pipe scope	[RFCXXXX]	
+-----+-----+-----+-----+-----+				

12.3. DOTS Signal Telemetry YANG Module

This document requests IANA to register the following URIs in the "ns" subregistry within the "IETF XML Registry" [RFC3688]:

URI: urn:ietf:params:xml:ns:yang:ietf-dots-telemetry
 Registrant Contact: The IESG.
 XML: N/A; the requested URI is an XML namespace.

URI: urn:ietf:params:xml:ns:yang:ietf-dots-mapping
 Registrant Contact: The IESG.
 XML: N/A; the requested URI is an XML namespace.

This document requests IANA to register the following YANG modules in the "YANG Module Names" subregistry [RFC6020] within the "YANG Parameters" registry.

name: ietf-dots-telemetry
 namespace: urn:ietf:params:xml:ns:yang:ietf-dots-telemetry
 maintained by IANA: N
 prefix: dots-telemetry
 reference: RFC XXXX

name: ietf-dots-mapping
 namespace: urn:ietf:params:xml:ns:yang:ietf-dots-mapping
 maintained by IANA: N
 prefix: dots-mapping
 reference: RFC XXXX

13. Security Considerations

13.1. DOTS Signal Channel Telemetry

The security considerations for the DOTS signal channel protocol are discussed in Section 12.11 of ~~[I-D.boucadair-dots-rfc8782-bis]~~. [\[I-D.ietf-dots-rfc8782-bis\]](#). The following discusses the security considerations that are specific to the DOTS signal channel extension defined in this document.

The DOTS telemetry information includes DOTS client network topology, DOTS client domain pipe capacity, normal traffic baseline and connections capacity, and threat and mitigation information. Such information is sensitive; it MUST be protected at rest by the DOTS server domain to prevent data leakage.

DOTS clients are typically trusted devices by the DOTS client domain. DOTS clients may be co-located on network security services (e.g., firewall) and a compromised security service potentially can do a lot more damage to the network. This assumption differs from the often held view that devices are untrusted, often referred to as the "zero-trust model". A compromised DOTS client can send fake DOTS telemetry data to a DOTS server to mislead the DOTS server. This attack can be prevented by monitoring and auditing DOTS clients to detect misbehavior and to deter misuse, and by only authorizing the DOTS client to convey the DOTS telemetry for specific target resources (e.g., an application server is authorized to exchange DOTS telemetry for its IP addresses but a DDoS mitigator can exchange DOTS telemetry for any target resource in the network). As a reminder, this is variation of dealing with compromised DOTS clients as discussed in Section 12.11 of ~~[I-D.boucadair-dots-rfc8782-bis]~~. [\[I-D.ietf-dots-rfc8782-bis\]](#).

DOTS servers must be capable of defending themselves against DoS attacks from compromised DOTS clients. The following non-comprehensive list of mitigation techniques can be used by a DOTS server to handle misbehaving DOTS clients:

- o The probing rate (defined in Section 5.5 4.5 of ~~[I-D.boucadair-dots-rfc8782-bis]~~ [\[I-D.ietf-dots-rfc8782-bis\]](#)) can be used to limit the average data rate to the DOTS server.

- o Rate-limiting DOTS telemetry, including those with new 'tmid' values, from the same DOTS client defends against DoS attacks that would result in varying the 'tmid' to exhaust DOTS server resources. Likewise, the DOTS server can enforce a quota and time-limit on the number of active pre-or-ongoing-mitigation telemetry data (identified by 'tmid') from the DOTS client.

Note also that telemetry notification interval may be used to rate-limit the pre-or-ongoing-mitigation telemetry notifications received by a DOTS client domain.

13.2. Vendor Attack Mapping

The security considerations for the DOTS data channel protocol are discussed in Section 10 of [RFC8783]. The following discusses the security considerations that are specific to the DOTS data channel extension defined in this document.

All data nodes defined in the YANG module specified in Section 10.2 which can be created, modified, and deleted (i.e., config true, which is the default) are considered sensitive. Write operations to these data nodes without proper protection can have a negative effect on network operations. Appropriate security measures are recommended to prevent illegitimate users from invoking DOTS data channel primitives as discussed in [RFC8783]. Nevertheless, an attacker who can access a DOTS client is technically capable of undertaking various attacks, such as:

- o Communicating invalid attack mapping details to the server
~~('/ietf-data:dots-data/ietf-data:dots-client/dots-~~
~~telemetry:vendor-mapping')~~, which will mislead the server when correlating attack details.

Some of the readable data nodes in the YANG module specified in Section 10.2 may be considered sensitive. It is thus important to control read access to these data nodes. These are the data nodes and their sensitivity:

- o ~~('/ietf-data:dots-data/ietf-data:dots-client/dots-telemetry:vendor-~~
~~mapping')~~ **('/data-channel:dots-data/data-channel:dots-client/dots-**
telemetry:vendor-mapping') can be misused to infer the DDoS protection technology deployed in a DOTS client domain.
- o ~~('/ietf-data:dots-data/dots-telemetry:vendor-mapping')~~ **('/data-channel:dots-data/dots-telemetry:vendor-mapping')** can be used by a compromised DOTS client to leak the attack detection capabilities of the DOTS server. This is a variation of the compromised DOTS client attacks discussed in Section 13.1.

14. Contributors

The following individuals have contributed to this document:

- o Li Su, CMCC, Email: suli@chinamobile.com
- o Pan Wei, Huawei, Email: william.panwei@huawei.com

15. Acknowledgements

The authors would like to thank Flemming Andreassen, Liang Xia, and Kaname Nishizuka co-authors of [I-D.doron-dots-telemetry] and everyone who had contributed to that document.

The authors would like to thank Kaname Nishizuka, Wei Pan, and Yuuhei Hayashi for comments and review.

Special thanks to Jon Shallow and Kaname Nishizuka for their implementation and interoperability work.

Many thanks to Jan Lindblad for the yangdoctors review and Nagendra Nainar for the opsdire review.

16. References

16.1. Normative References

[Enterprise-Numbers]
 "Private Enterprise Numbers", May 2020,
<http://www.iana.org/assignments/enterprise-numbers.html>.

~~[I-D.boucadair-dots-rfc8782-bis]~~

[I-D.ietf-dots-rfc8782-bis]

Boucadair, M., ~~Ed.~~, Shallow, J., and T. Reddy.K, "Distributed Denial-of-Service Open Threat Signaling (DOTS) Signal Channel Specification",
<https://tools.ietf.org/html/draft-boucadair-dots-rfc8782-bis-00>, **draft-ietf-dots-rfc8782-bis-01**

(work in progress), September 2020.

- [I-D.ietf-dots-signal-filter-control]
Nishizuka, K., Boucadair, M., Reddy, K. T., and T. Nagata,
"Controlling Filtering Rules Using Distributed Denial-of-
Service Open Threat Signaling (DOTS) Signal Channel",
draft-ietf-dots-signal-filter-control-07 (work in
progress), June 2020.
- [RFC2119] Bradner, S., "Key words for use in RFCs to Indicate
Requirement Levels", BCP 14, RFC 2119,
DOI 10.17487/RFC2119, March 1997,
<<https://www.rfc-editor.org/info/rfc2119>>.
- [RFC3688] Mealling, M., "The IETF XML Registry", BCP 81, RFC 3688,
DOI 10.17487/RFC3688, January 2004,
<<https://www.rfc-editor.org/info/rfc3688>>.
- [RFC6020] Bjorklund, M., Ed., "YANG - A Data Modeling Language for
the Network Configuration Protocol (NETCONF)", RFC 6020,
DOI 10.17487/RFC6020, October 2010,
<<https://www.rfc-editor.org/info/rfc6020>>.
- [RFC6991] Schoenwaelder, J., Ed., "Common YANG Data Types",
RFC 6991, DOI 10.17487/RFC6991, July 2013,
<<https://www.rfc-editor.org/info/rfc6991>>.
- [RFC7049] Bormann, C. and P. Hoffman, "Concise Binary Object
Representation (CBOR)", RFC 7049, DOI 10.17487/RFC7049,
October 2013, <<https://www.rfc-editor.org/info/rfc7049>>.
- [RFC7252] Shelby, Z., Hartke, K., and C. Bormann, "The Constrained
Application Protocol (CoAP)", RFC 7252,
DOI 10.17487/RFC7252, June 2014,
<<https://www.rfc-editor.org/info/rfc7252>>.
- [RFC7641] Hartke, K., "Observing Resources in the Constrained
Application Protocol (CoAP)", RFC 7641,
DOI 10.17487/RFC7641, September 2015,
<<https://www.rfc-editor.org/info/rfc7641>>.
- [RFC7950] Bjorklund, M., Ed., "The YANG 1.1 Data Modeling Language",
RFC 7950, DOI 10.17487/RFC7950, August 2016,
<<https://www.rfc-editor.org/info/rfc7950>>.
- [RFC7959] Bormann, C. and Z. Shelby, Ed., "Block-Wise Transfers in
the Constrained Application Protocol (CoAP)", RFC 7959,
DOI 10.17487/RFC7959, August 2016,
<<https://www.rfc-editor.org/info/rfc7959>>.
- [RFC7970] Danyliw, R., "The Incident Object Description Exchange
Format Version 2", RFC 7970, DOI 10.17487/RFC7970,
November 2016, <<https://www.rfc-editor.org/info/rfc7970>>.
- [RFC8040] Bierman, A., Bjorklund, M., and K. Watsen, "RESTCONF
Protocol", RFC 8040, DOI 10.17487/RFC8040, January 2017,
<<https://www.rfc-editor.org/info/rfc8040>>.
- [RFC8174] Leiba, B., "Ambiguity of Uppercase vs Lowercase in RFC
2119 Key Words", BCP 14, RFC 8174, DOI 10.17487/RFC8174,
May 2017, <<https://www.rfc-editor.org/info/rfc8174>>.
- [RFC8345] Clemm, A., Medved, J., Varga, R., Bahadur, N.,
Ananthakrishnan, H., and X. Liu, "A YANG Data Model for
Network Topologies", RFC 8345, DOI 10.17487/RFC8345, March
2018, <<https://www.rfc-editor.org/info/rfc8345>>.
- [RFC8768] Boucadair, M., Reddy, K. T., and J. Shallow, "Constrained
Application Protocol (CoAP) Hop-Limit Option", RFC 8768,
DOI 10.17487/RFC8768, March 2020,
<<https://www.rfc-editor.org/info/rfc8768>>.
- [RFC8783] Boucadair, M., Ed. and T. Reddy, K., Ed., "Distributed
Denial-of-Service Open Threat Signaling (DOTS) Data
Channel Specification", RFC 8783, DOI 10.17487/RFC8783,
May 2020, <<https://www.rfc-editor.org/info/rfc8783>>.
- [RFC8791] Bierman, A., Bjorklund, M., and K. Watsen, "YANG Data
Structure Extensions", RFC 8791, DOI 10.17487/RFC8791,
June 2020, <<https://www.rfc-editor.org/info/rfc8791>>.

16.2. Informative References

- [Cause] IANA, "DOTS Signal Channel Conflict Cause Codes",
<<https://www.iana.org/assignments/dots/dots.xhtml#dots-signal-channel-conflict-cause-codes>>.

~~{I-D.bosh-core-new-block}~~

~~Boucadair, M. and J. Shallow, "Constrained Application Protocol (CoAP) Block-Wise Transfer Options for Faster Transmission", draft-bosh-core-new-block-04 (work in progress), June 2020.~~

[I-D.doron-dots-telemetry]
Doron, E., Reddy, T., Andreasen, F., Xia, L., and K. Nishizuka, "Distributed Denial-of-Service Open Threat Signaling (DOTS) Telemetry Specifications", draft-doron-dots-telemetry-00 (work in progress), October 2016.

[I-D.ietf-core-new-block]
Boucadair, M. and J. Shallow, "Constrained Application Protocol (CoAP) Block-Wise Transfer Options for Faster Transmission", draft-ietf-core-new-block-01 (work in progress), September 2020.

[I-D.ietf-dots-multihoming]
Boucadair, M., Reddy, K. T., and W. Pan, "Multi-homing Deployment Considerations for Distributed-Denial-of-Service Open Threat Signaling (DOTS)", draft-ietf-dots-multihoming-04 (work in progress), May 2020.

[I-D.ietf-dots-use-cases]
Dobbins, R., Migault, D., Moskowitz, R., Teague, N., Xia, L., and K. Nishizuka, "Use cases for DDoS Open Threat Signaling", draft-ietf-dots-use-cases-25 (work in progress), July 2020.

[Key-Map] IANA, "DOTS Signal Channel CBOR Key Values", <<https://www.iana.org/assignments/dots/dots.xhtml#dots-signal-channel-cbor-key-values>>.

[RFC2330] Paxson, V., Almes, G., Mahdavi, J., and M. Mathis, "Framework for IP Performance Metrics", RFC 2330, DOI 10.17487/RFC2330, May 1998, <<https://www.rfc-editor.org/info/rfc2330>>.

[RFC8340] Bjorklund, M. and L. Berger, Ed., "YANG Tree Diagrams", BCP 215, RFC 8340, DOI 10.17487/RFC8340, March 2018, <<https://www.rfc-editor.org/info/rfc8340>>.

[RFC8612] Mortensen, A., Reddy, T., and R. Moskowitz, "DDoS Open Threat Signaling (DOTS) Requirements", RFC 8612, DOI 10.17487/RFC8612, May 2019, <<https://www.rfc-editor.org/info/rfc8612>>.

Authors' Addresses

Mohamed Boucadair (editor)
Orange
Rennes 35000
France

Email: mohamed.boucadair@orange.com

Tirumaleswar Reddy (editor)
McAfee, Inc.
Embassy Golf Link Business Park
Bangalore, Karnataka 560071
India

Email: kondtir@gmail.com
Ehud Doron
Radware Ltd.
Raoul Wallenberg Street
Tel-Aviv 69710
Israel

Email: ehudd@radware.com

Meiling Chen
CMCC
32, Xuanwumen West
BeiJing, BeiJing 100053
China

Email: chenmeiling@chinamobile.com

Jon Shallow
United Kingdom

Email: supjps-ietf@jpshallow.com